

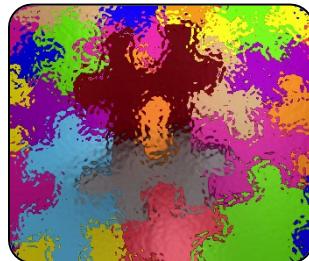
Introduction to Cloud COMPUTING

Giuliano Taffoni - I.N.A.F.

“Cloud Service Model”
Lecture 4

2023 @ Università di Trieste

Outline



Intro to Cloud Service
Models

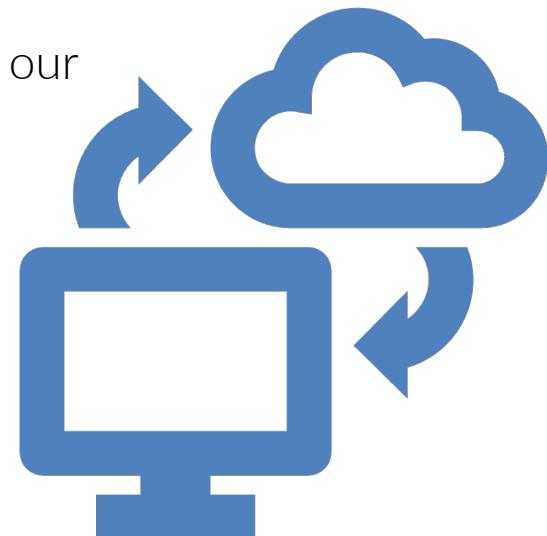


IaaS & PaaS

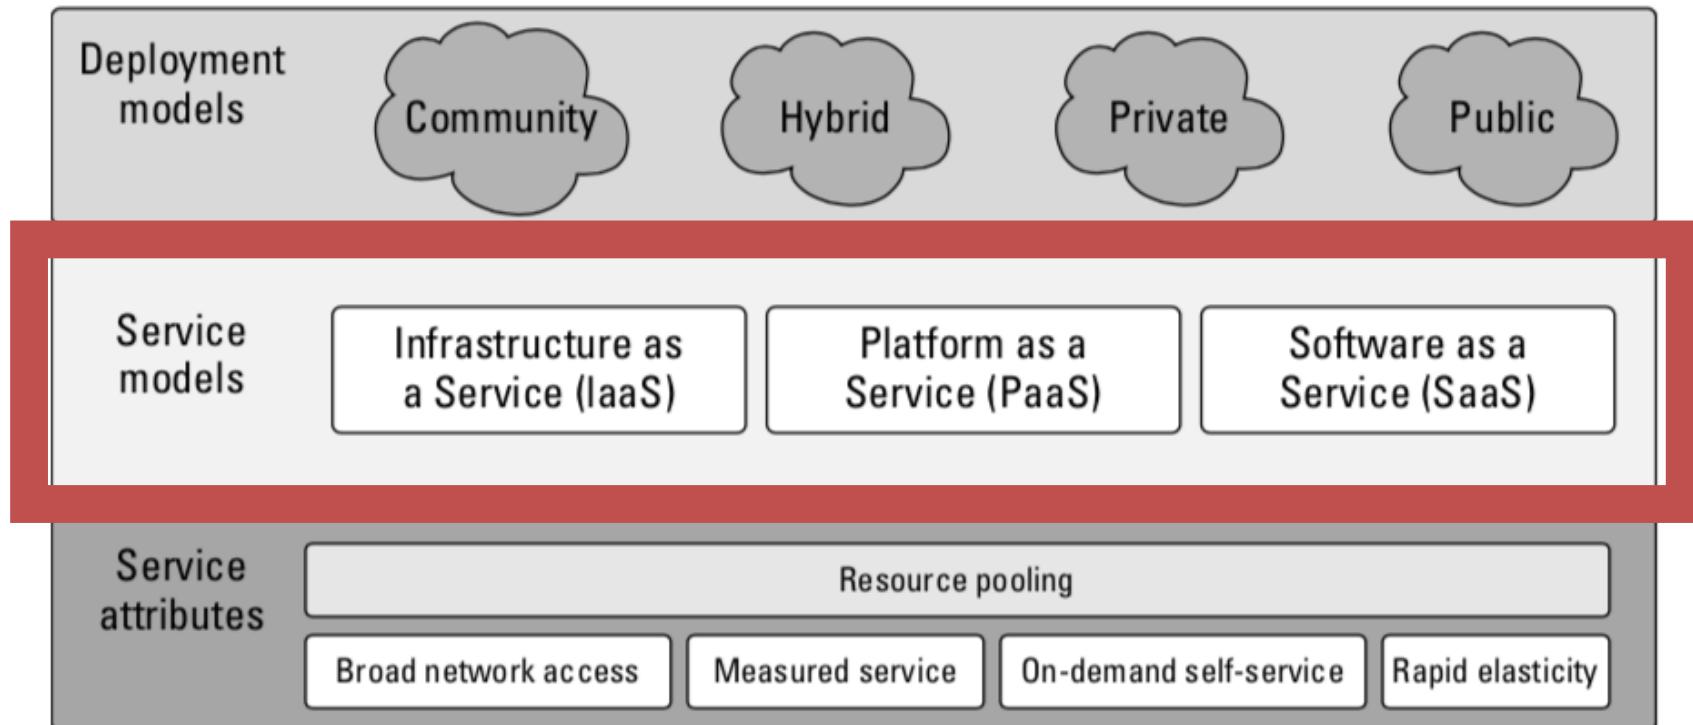
| Cloud Definition: again



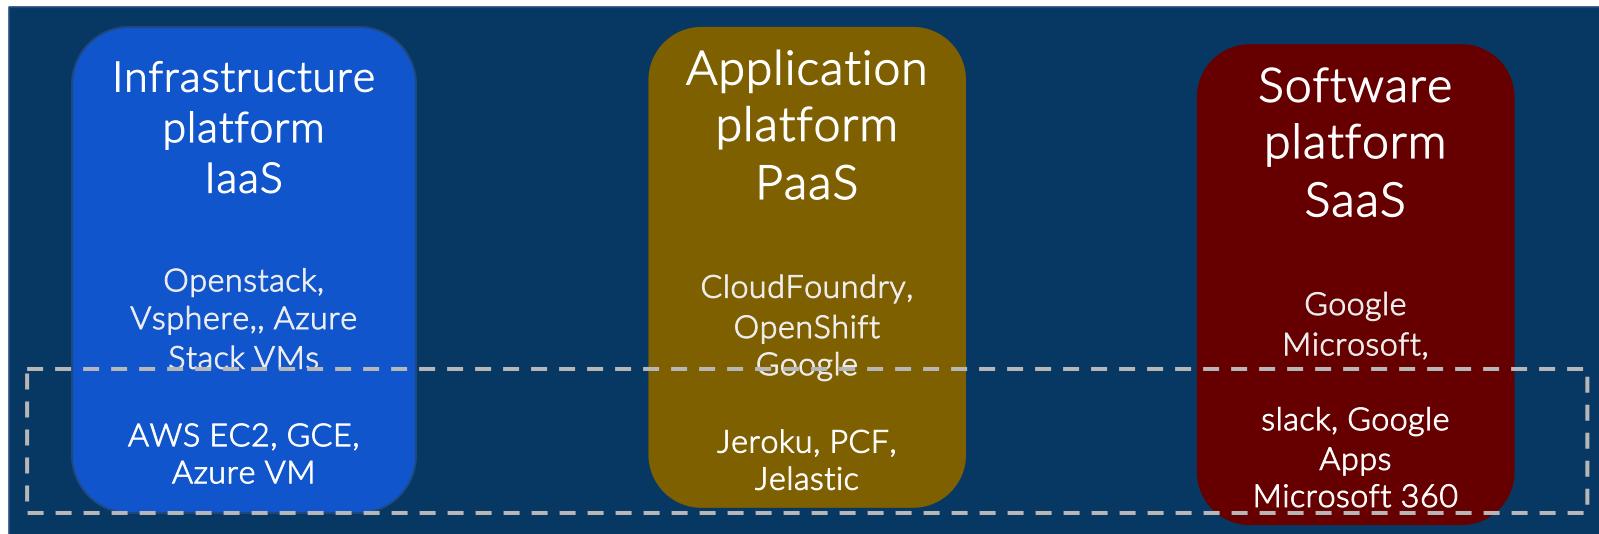
- I. A model of computation and data storage based on “pay as you go”, access to “unlimited” remote data center capabilities
- II. A cloud infrastructure provides a framework to manage scalable, reliable, on-demand access to applications
- III. Cloud services provide the “invisible” backend to many of our mobile applications
- IV. High level of elasticity in consumption
- V. Historical roots in today’s Internet apps
 - Search, email, social networks
 - File storage (Live Mesh, Mobile Me, Flickr, ...)



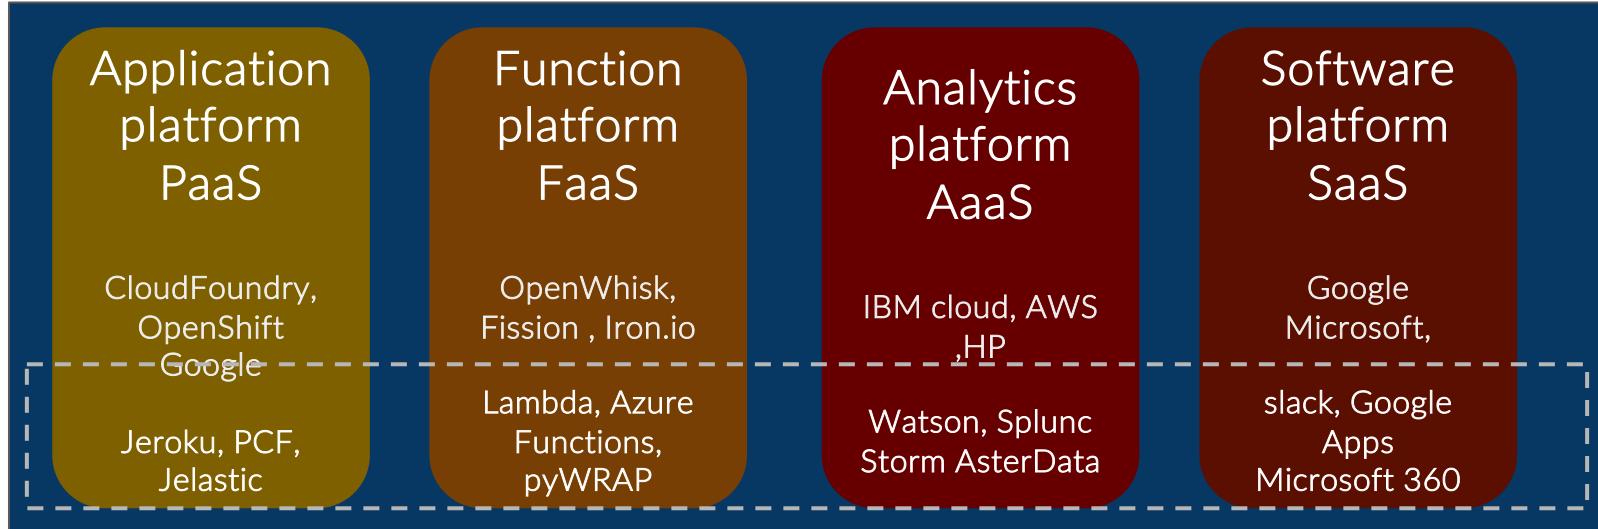
Cloud Architecture



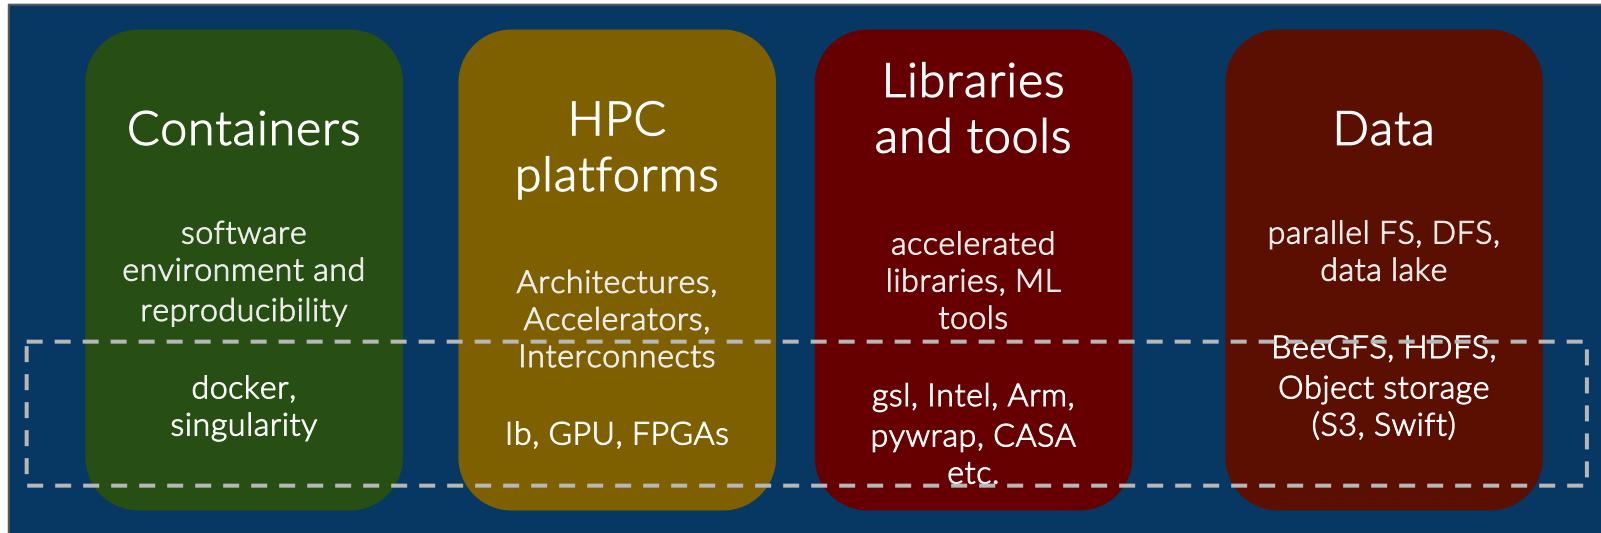
Traditional service models



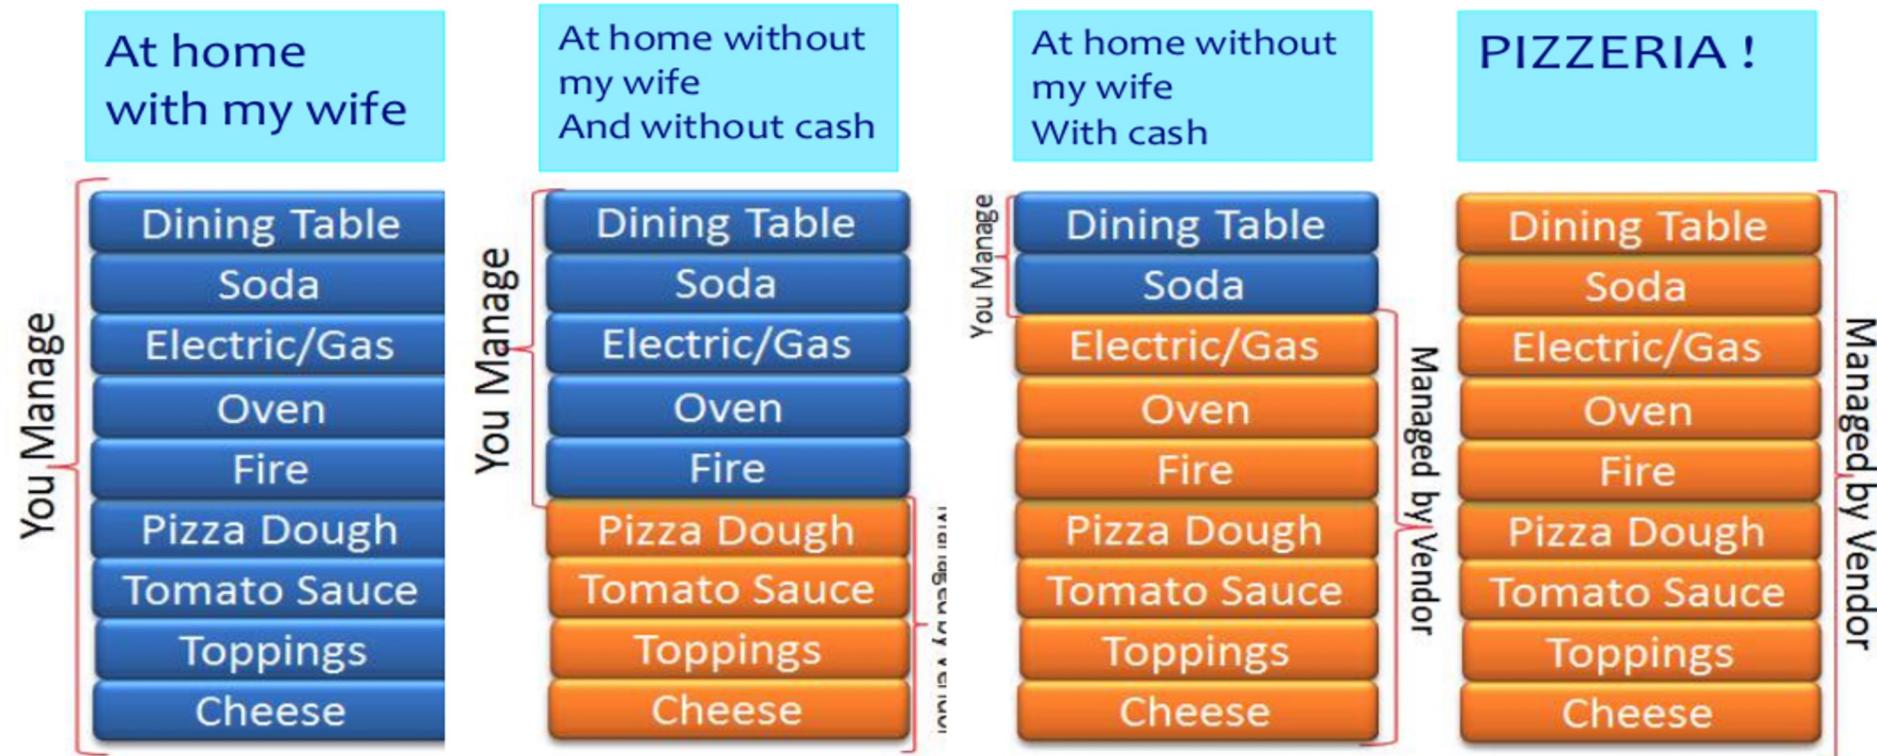
Modern Models



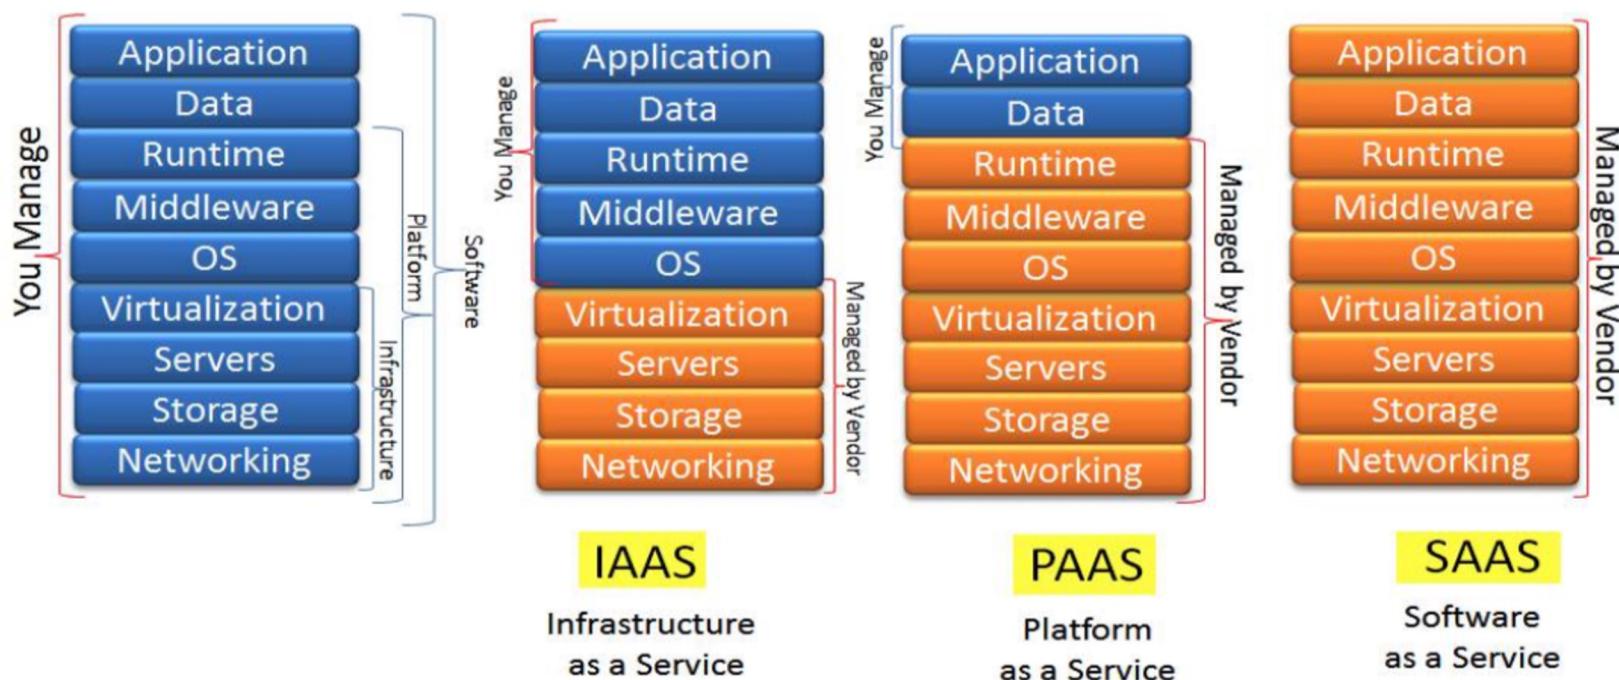
Modern Models



When I want a Pizza



When I manage IT stuff...



Taken and Slightly changed from

<https://cloudcelebrity.wordpress.com/2011/11/22/introduction-to-cloud-services-iaas-paas-saas/>

What is cloud?

A cloud infrastructure is:

An *infrastructure* to provide users with the most **flexible** way to allocate computational power and storage space (and any other IT services)

To quickly provide a customized IT infrastructure

...and you must integrate infrastructure provisioning in your data-analysis.

| For example in OpenStack....

```
>>> from keystoneauth1 import loading
>>> from keystoneauth1 import session
>>> from novaclient import client
>>> loader = loading.get_plugin_loader('password')
>>> auth =
loader.load_from_options(auth_url=AUTH_URL,
...
username=USERNAME,
...
password=PASSWORD,
...
project_id=PROJECT_ID)
>>> sess = session.Session(auth=auth)
>>> nova = client.Client(VERSION, session=sess)
```

| Infrastructure as a Service

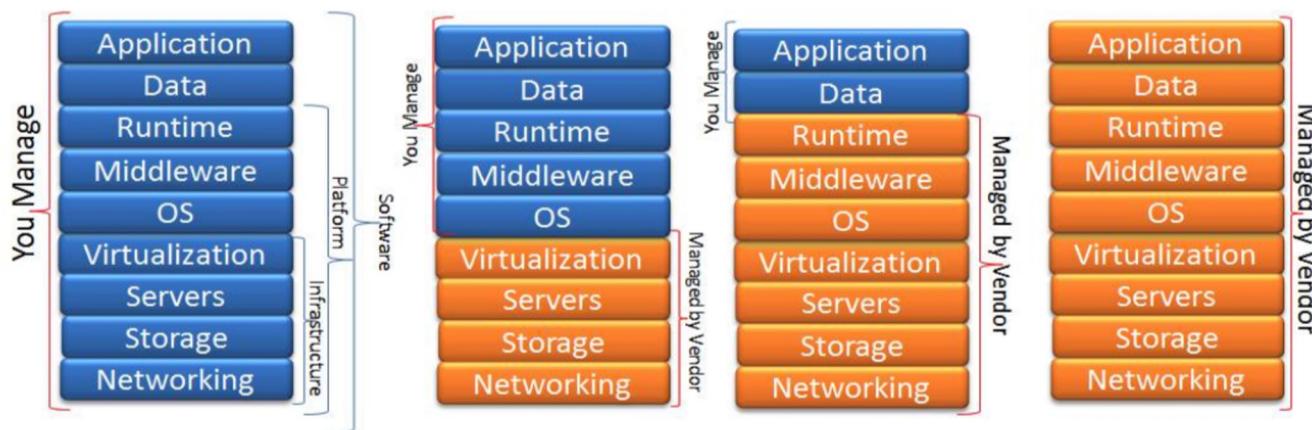
"Infrastructure as a service (IaaS) refers to online services that provide high-level APIs used to dereference various low-level details of underlying network infrastructure like physical computing resources, location, data partitioning, scaling, security, backup etc. A **hypervisor**, such as Xen, Oracle VirtualBox, Oracle VM, KVM, VMware ESX/ESXi, or Hyper-V, LXD, runs the virtual machines as guests. Pools of hypervisors within the **cloud operational system** can support large numbers of virtual machines and the ability to **scale services up and down** according to customers' varying requirements."

Wikipedia

IaaS

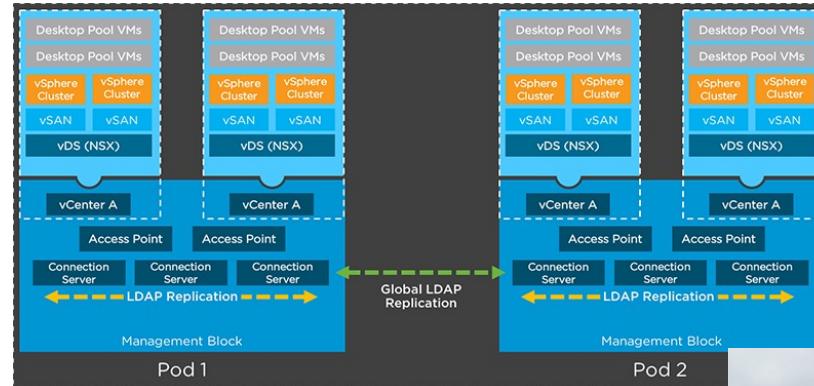
Infrastructure as a Service is a cloud computing service model in which hardware is virtualized in the cloud.

Separation of responsibilities



IaaS architectural elements

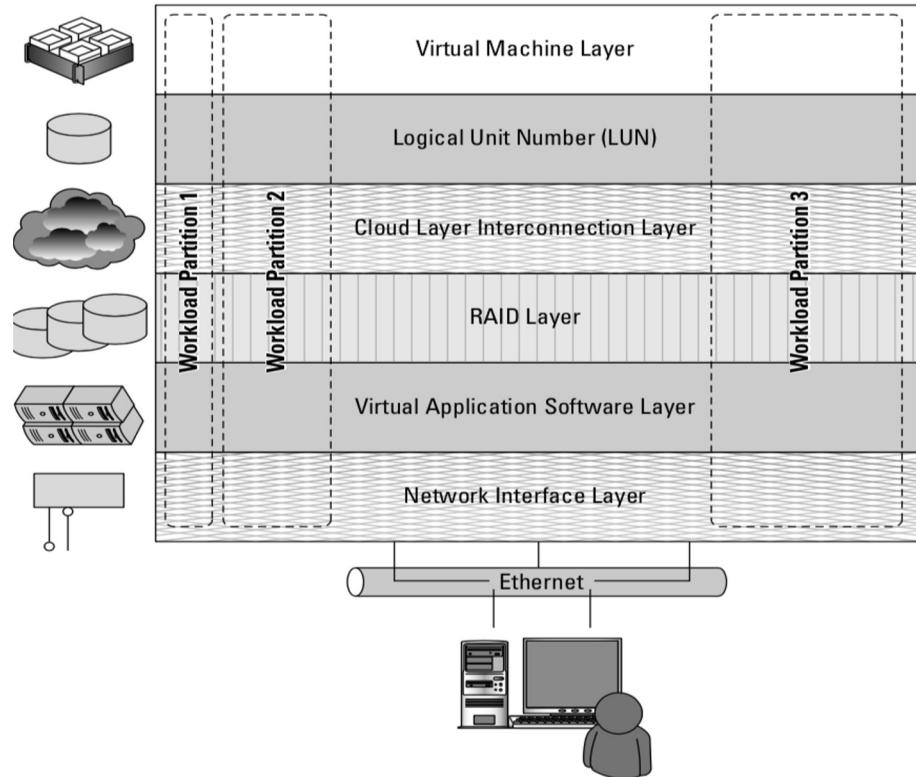
- Workloads
- Point of Delivery
- Aggregation
- Silos



IaaS Workloads

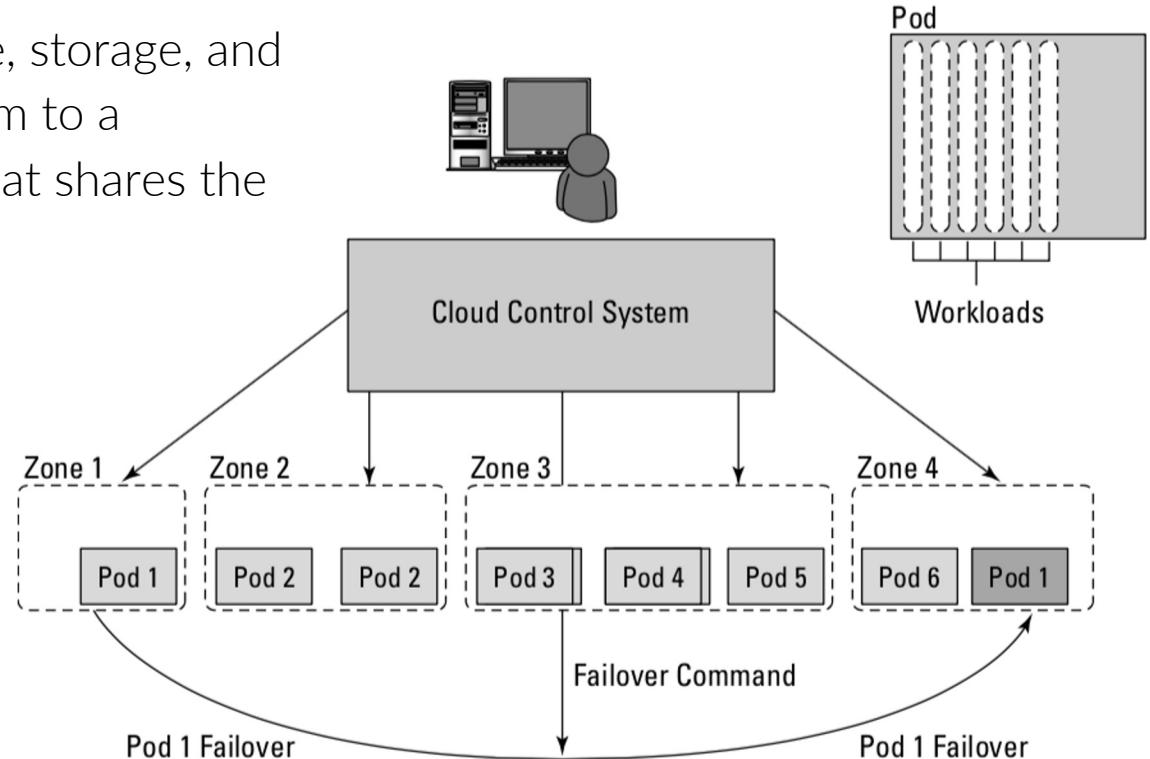
A cloud workload is a specific application, service, capability or a specific amount of work that can be run on a cloud resource. Virtual machines, databases, containers,

Hadoop nodes and applications are all considered cloud workloads.



Point of Delivery

POD is a collection of compute, storage, and network resources that conform to a standard operating footprint that shares the same failure domain.



POD vs SILO

POD are aggregated into pools within the same IaaS zone or area: the *availability zone*.

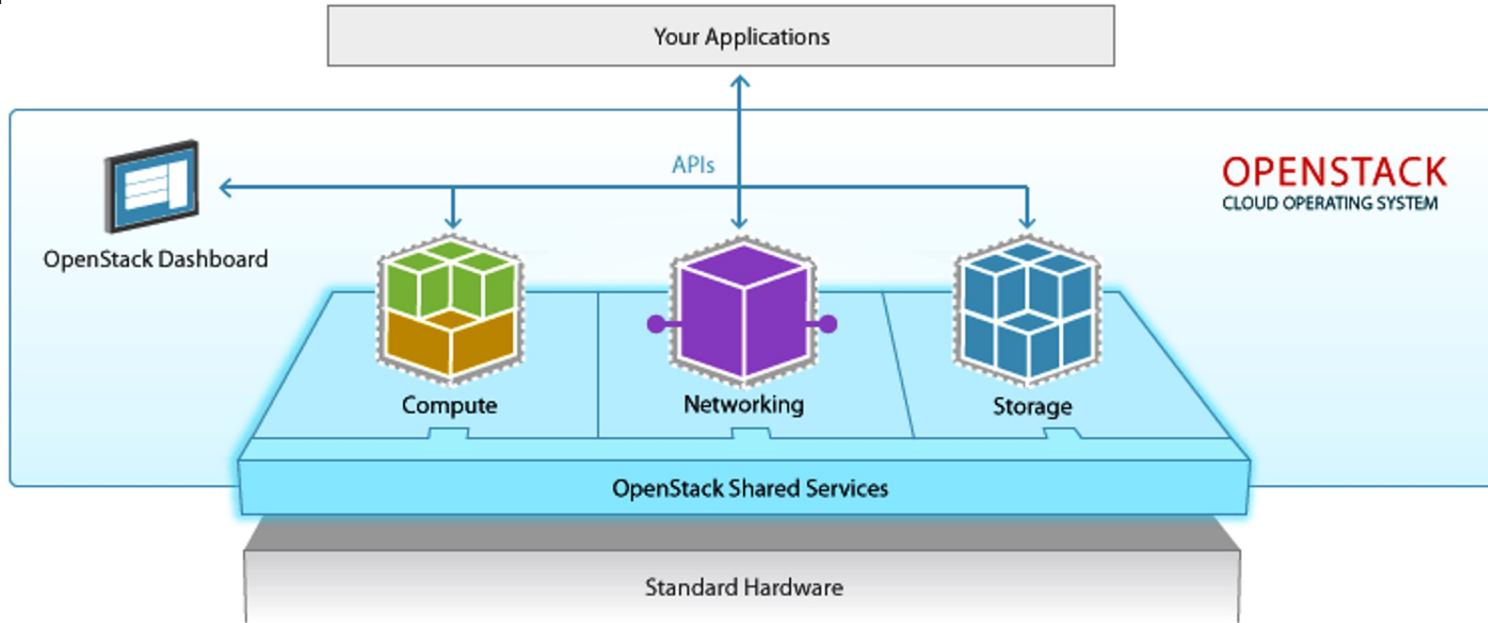
When cloud computing infrastructure isolates user clouds from each other so the management system is incapable of interoperating with other private clouds, it creates an information silo.

IaaS examples

Amazon Web Services (AWS)
Google Compute Engine (GCE)
IBM Cloud
Microsoft Azure
Rackspace
Linode
Cisco Metacloud
Digital Ocean
Vultr
Oracle Cloud

European Open Science Cloud (OpenStack) (<https://eosc-portal.eu/>)

OpenStack



OpenStack is a **cloud operating system** that controls large **pools** of compute, storage, and networking resources throughout a datacenter, all managed through a **dashboard** that gives administrators control while **empowering their users** to provision resources through a web Interface.

OpenStack main concepts

- The OpenStack project is an open source cloud computing platform that supports all types of cloud environments.
- The project aims for simple implementation, massive scalability, and a rich set of features.
- Cloud computing experts from around the world contribute to the project.
- Based on services (composability), each service offers an application programming interface (API) that facilitates this integration

OpenStack principles

- A. **Open development model:** Code reviews and roadmaps are public. The code for OpenStack is freely available under the Apache 2.0 license.
- B. **Open design process:** Every six months the development community holds a design summit to gather requirements and write specifications for the upcoming release. The summit is open to the public and attendees include users, operators, developers, and upstream project personnel.
- C. **Open community:** OpenStack is dedicated to producing a healthy, vibrant, and active developer and user community. Most decisions are made by consensus. All processes are documented, open and transparent.

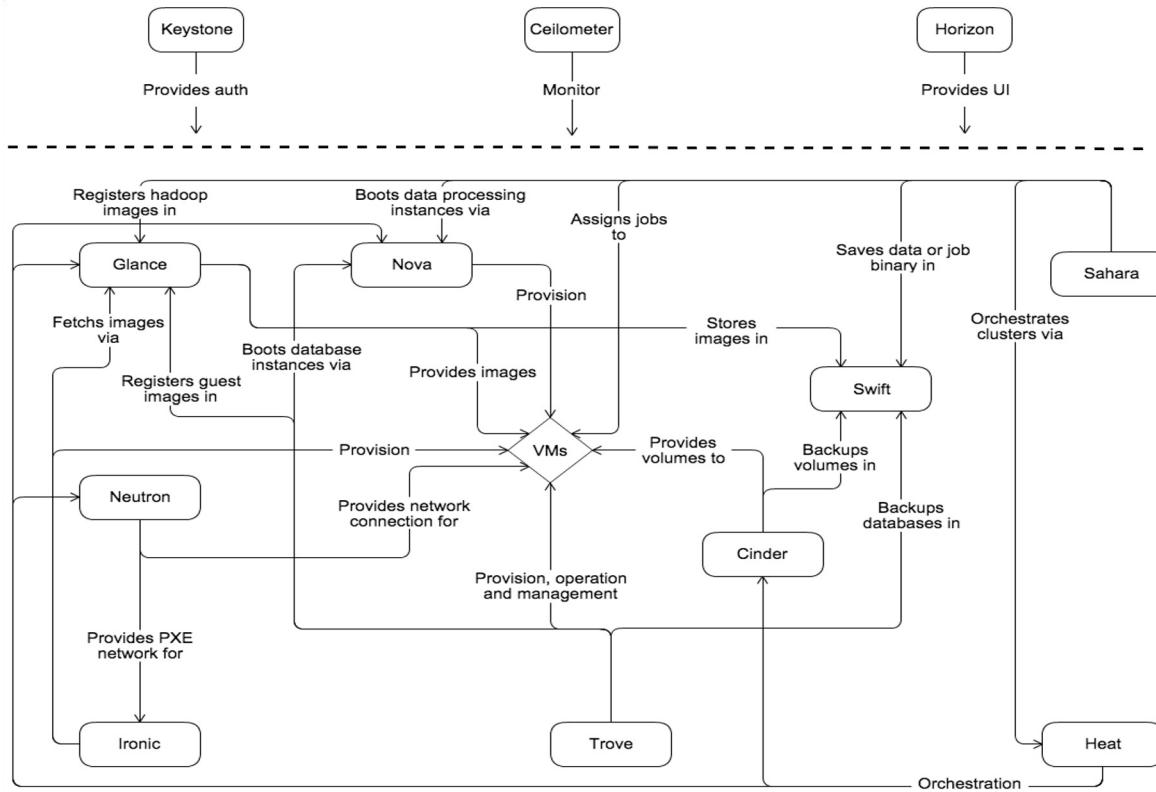
| OpenStack Capabilities

- Virtual machines (VMs) and Containers on demand
 - provisioning
 - snapshotting
- Networks
- Volumes
- Storage for VMs and arbitrary files
- Multi-tenancy
 - quotas for different projects, users
 - user can be associated with multiple projects

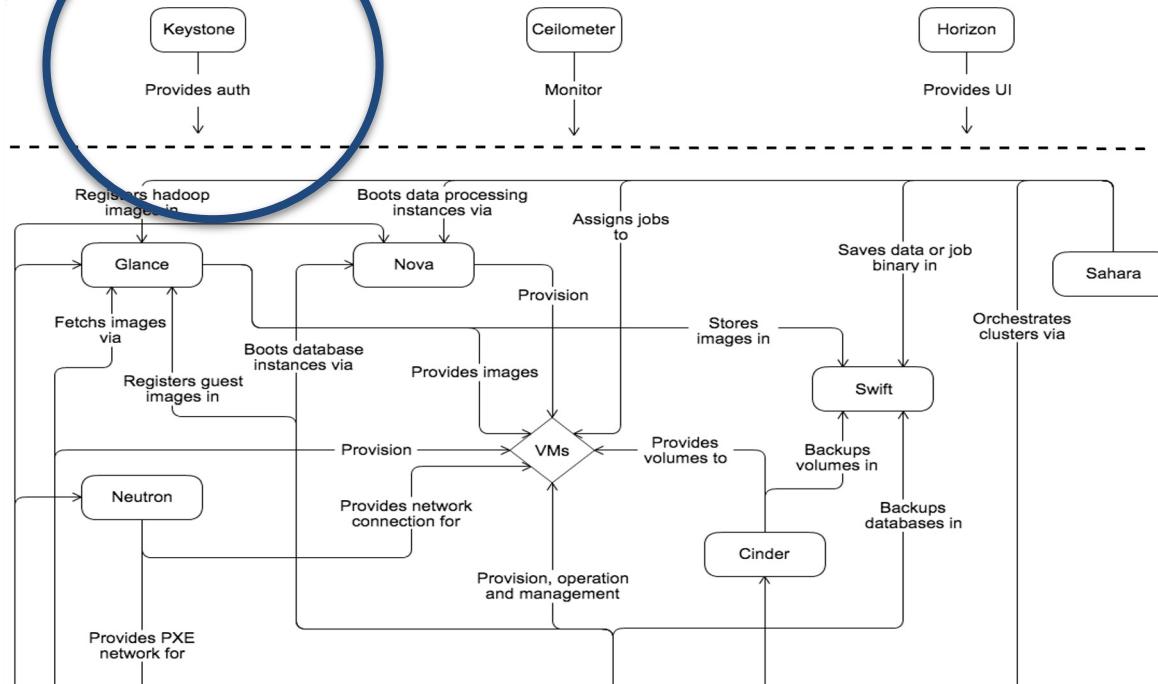
OpenStack tech details

- Written in Python (plus auxiliary shell scripts)
- Built around independent components
- Highly distributed architecture
 - designed for very big installations
- Intrinsic HA of most OpenStack services
 - MySQL and RabbitMQ have to be properly configured
- *SQL database used to store persistent data
- RabbitMQ used for RPC and notification
- RESTful APIs for all the services

OpenStack IaaS services view



OpenStack IaaS services view

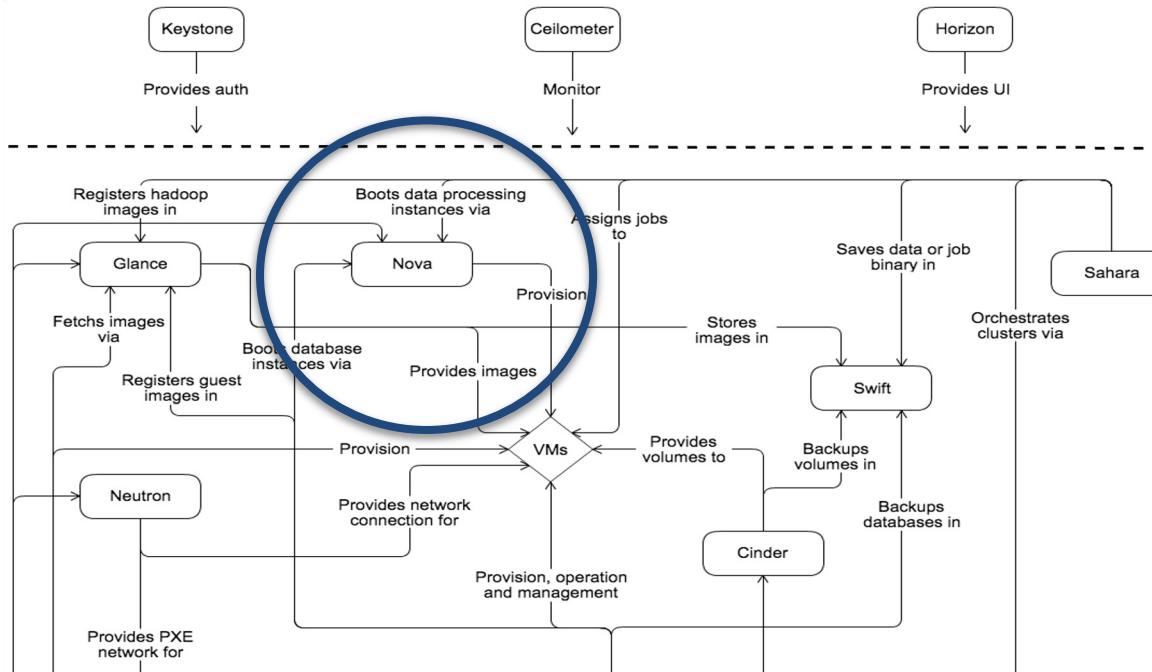


Keystone:authentication service

Keystone

```
>>> fromkeystoneauth1.identityimportv3
>>> fromkeystoneauth1importsession
>>> fromkeystoneclient.v3importclient
>>> auth=v3.Password(auth_url='https://my.keystone.com:5000/v3',
...         user_id='myuserid',
...         password='mypassword',
...         project_id='myprojectid')
>>> sess=session.Session(auth=auth)
>>> keystone=client.Client(session=sess)
```

OpenStack IaaS services view



Nova: computational service

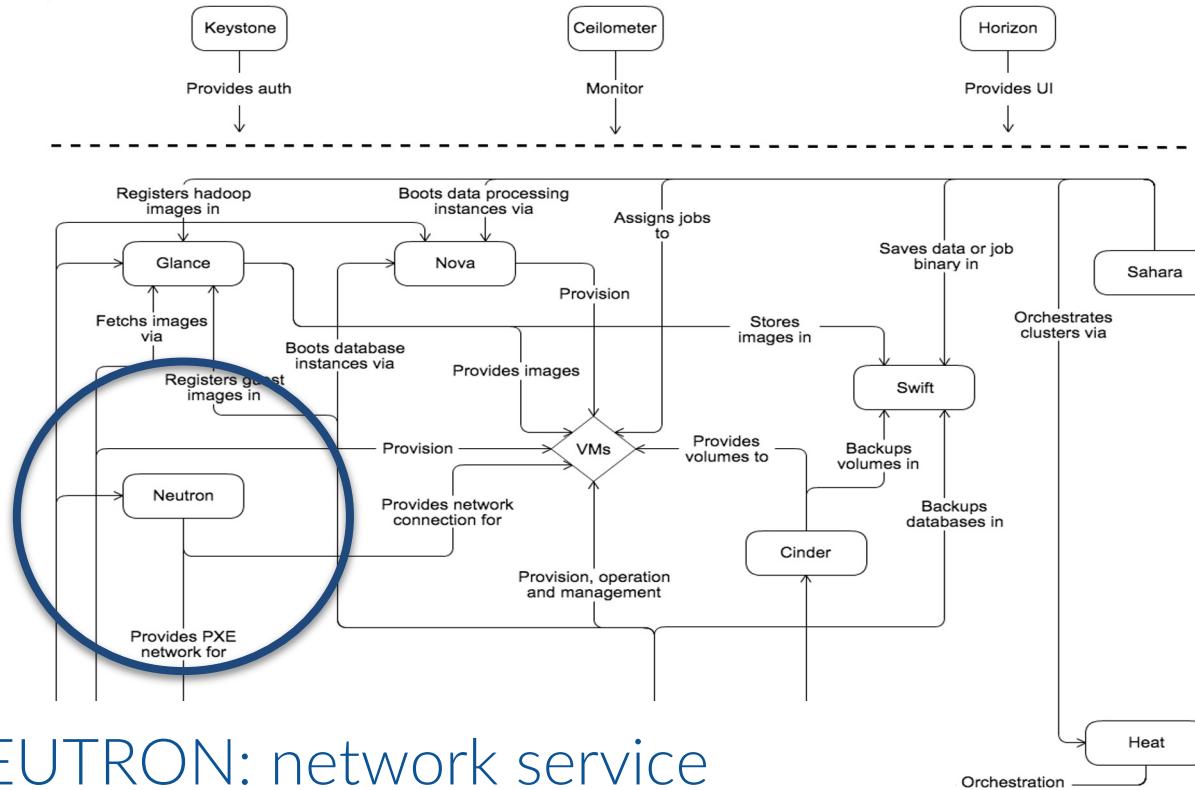
```
>>> nova = client.Client(VERSION, session=sess)

...
>>> nova.servers.list()
[<Server: buildslave-ubuntu-9.10>]

>>> nova.flavors.list()
[<Flavor: 256 server>,
 <Flavor: 512 server>,
 <Flavor: 1GiB server>,
 <Flavor: 2GiB server>,
 <Flavor: 4GiB server>,
 <Flavor: 8GiB server>,
 <Flavor: 15.5GiB server>]

>>> fl = nova.flavors.find(ram=512)
>>> nova.servers.create("my-server", flavor=fl)
<Server: my-server>
```

OpenStack IaaS services view



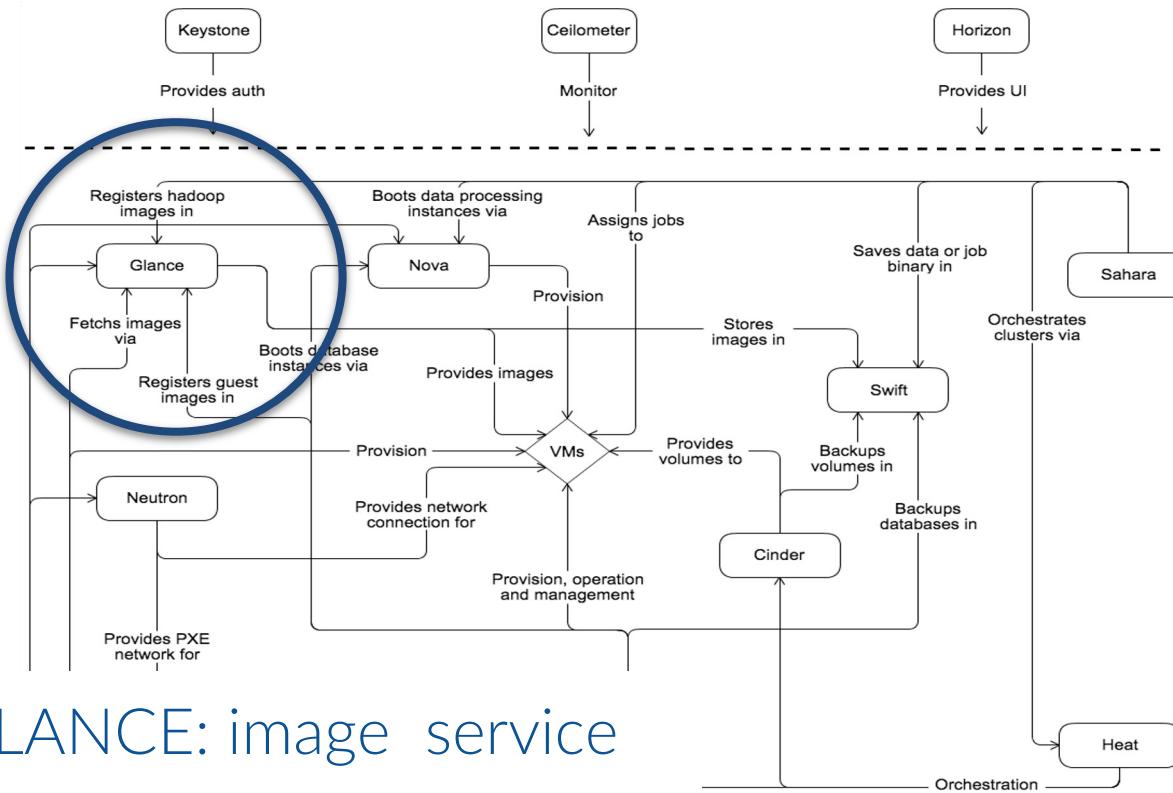
NEUTRON: network service

```
>>> neutron = client.Client(session=sess)

...
...

>>> network = {'name': 'mynetwork', 'admin_state_up': True}
>>> neutron.create_network({'network':network})
>>> networks = neutron.list_networks(name='mynetwork')
>>> print networks
>>> network_id = networks['networks'][0]['id']
>>> neutron.delete_network(network_id)
```

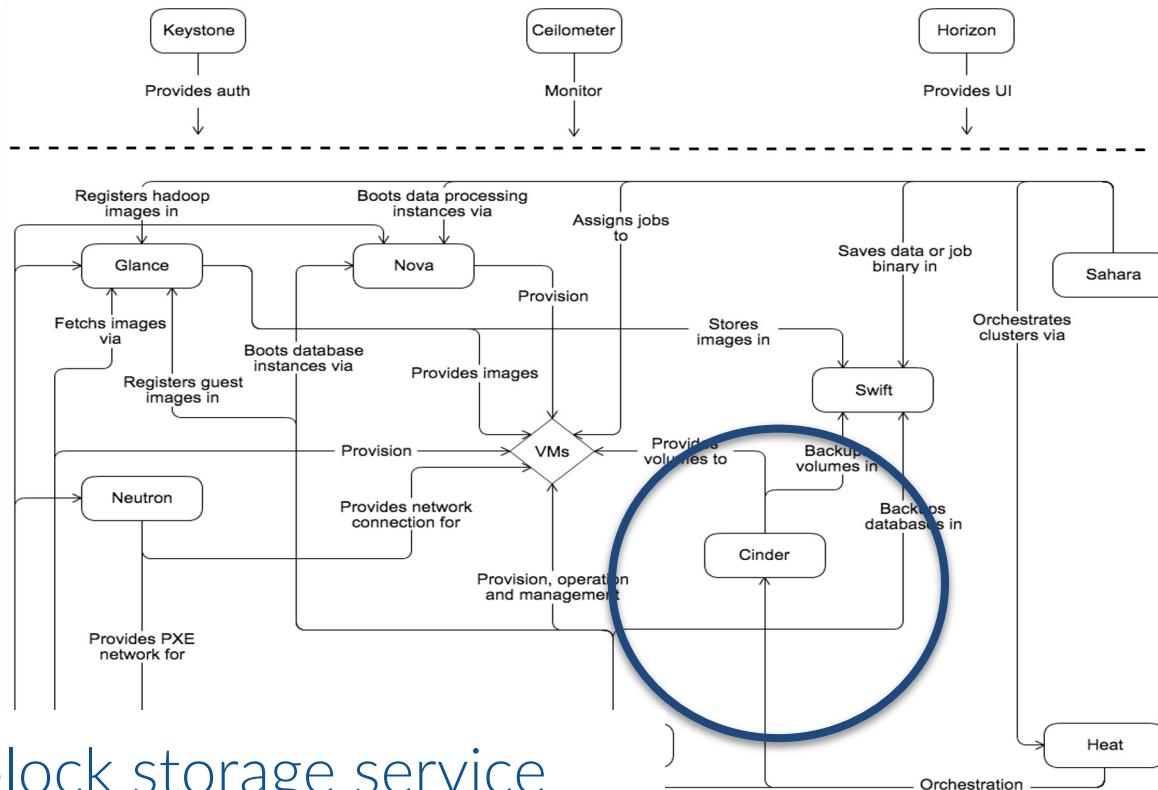
OpenStack IaaS services view



GLANCE: image service

```
>>> from glanceclient import Client  
  
>>> glance = Client('1', endpoint=OS_IMAGE_ENDPOINT, token=OS_AUTH_TOKEN)  
  
>>> image = glance.images.create(name="My Test Image")  
  
>>> print image.status  
  
'queued'  
  
>>> image.update(data=open('/tmp/myimage.iso', 'rb'))  
  
>>> print image.status  
  
'active'  
  
>>> image.update(properties=dict(my_custom_property='value'))  
  
>>> with open('/tmp/copyimage.iso', 'wb') as f:  
  
    for chunk in image.data():  
  
        f.write(chunk)  
  
>>> image.delete()
```

OpenStack IaaS services view



CINDER: block storage service

```
>>> sess = session.Session(auth=auth)

>>> cinder = client.Client(VERSION, session=sess)

>>> cinder.volumes.list()

>>> myvol = cinder.volumes.create(display_name="test-vol", size=1)

>>> myvol.id

ce06d0a8-5c1b-4e2c-81d2-39eca6bbfb70

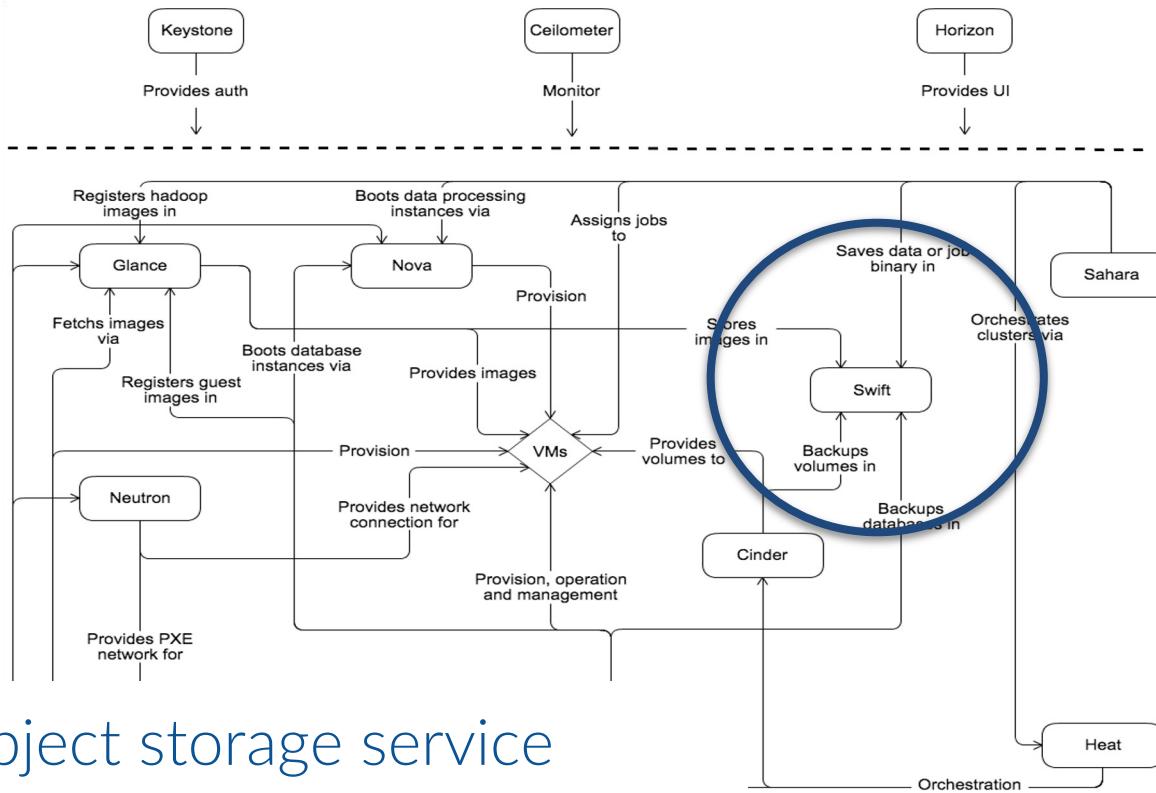
>>> cinder.volumes.list()

[<Volume: ce06d0a8-5c1b-4e2c-81d2-39eca6bbfb70>]

>>>myvol.delete

[]
```

OpenStack IaaS services view



SWIFT: object storage service

```
container = 'new-container'

conn.put_container(container)

resp_headers, containers = conn.get_account()

if container in containers:

    print("The container was created")


container = 'new-container'

with open('local.txt', 'r') as local:

    conn.put_object(

        container,
        'local_object.txt',
        contents=local,
        content_type='text/plain'

    )

obj = 'local_object.txt'

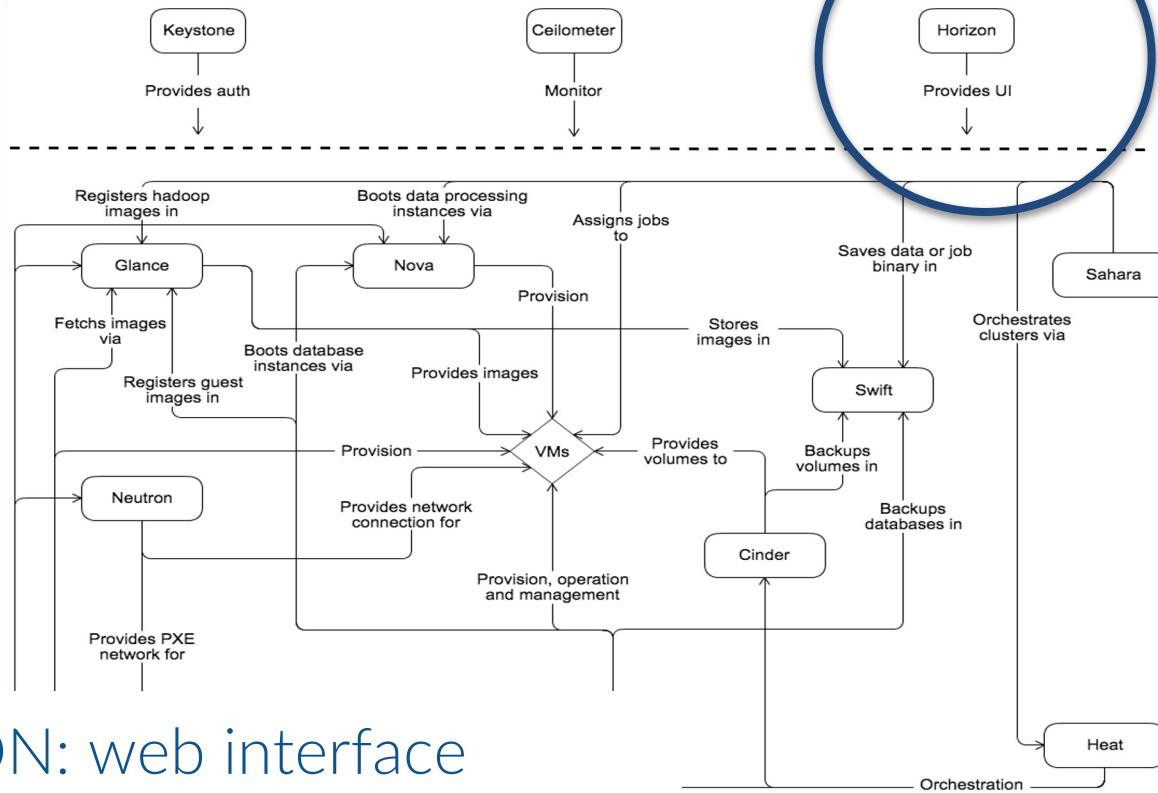
container = 'new-container'

resp_headers, obj_contents = conn.get_object(container, obj)

with open('local_copy.txt', 'w') as local:

    local.write(obj_contents)
```

OpenStack IaaS services view



HORIZON: web interface

Project

CURRENT PROJECT
sct

Manage Compute

Overview

Instances

Volumes

Images & Snapshots

Access & Security

Manage Network

Network Topology

Networks

Routers

Object Store

Containers

Overview

Limit Summary



Instances
Used 11 of 60



VCPUs
Used 58 of 200



RAM
Used 58.0 GB of 4.9 TB



Floating IPs
Used 57 of 60



Security Groups
Used 0 of 10

Select a period of time to query its usage:

From: 2014-07-01

To: 2014-07-11

Submit

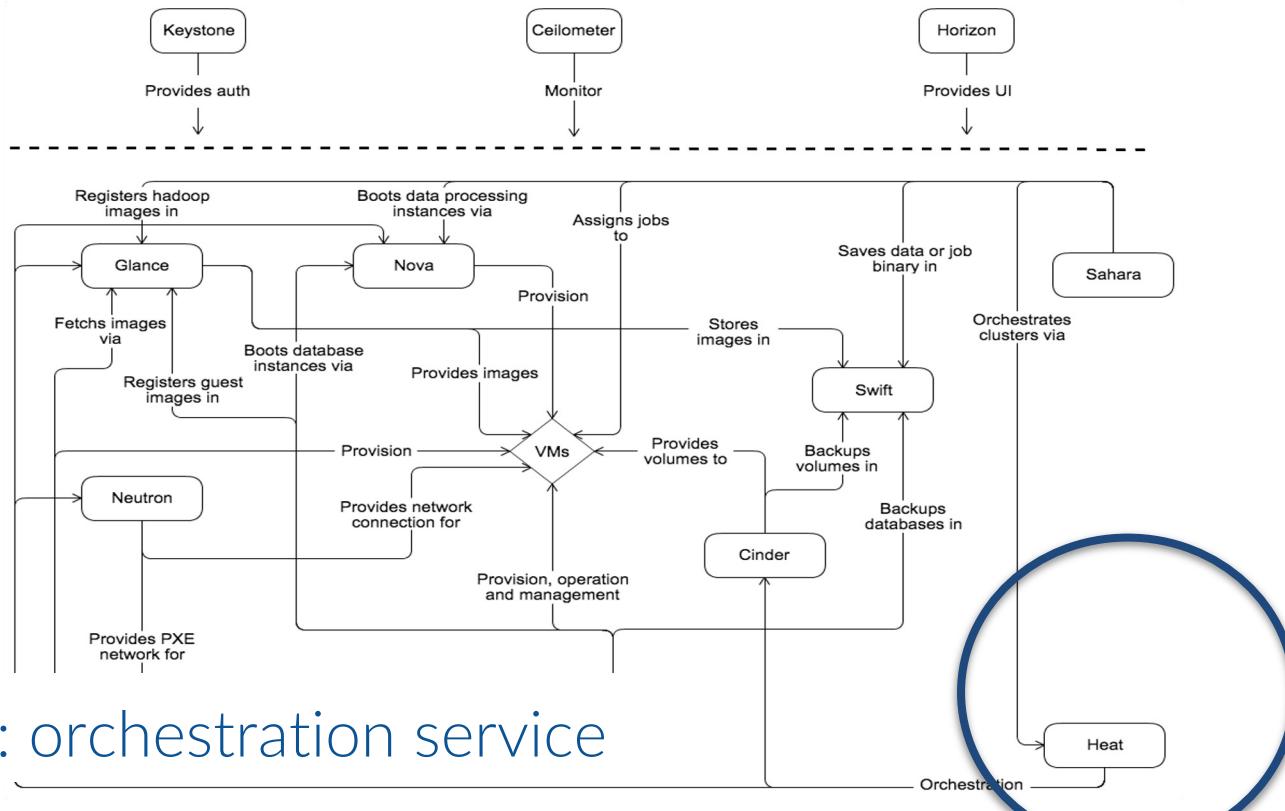
The date should be in YYYY-mm-dd format.

Active Instances: 11 Active RAM: 58GB This Period's VCPU-Hours: 9.86 This Period's GB-Hours: 215.17

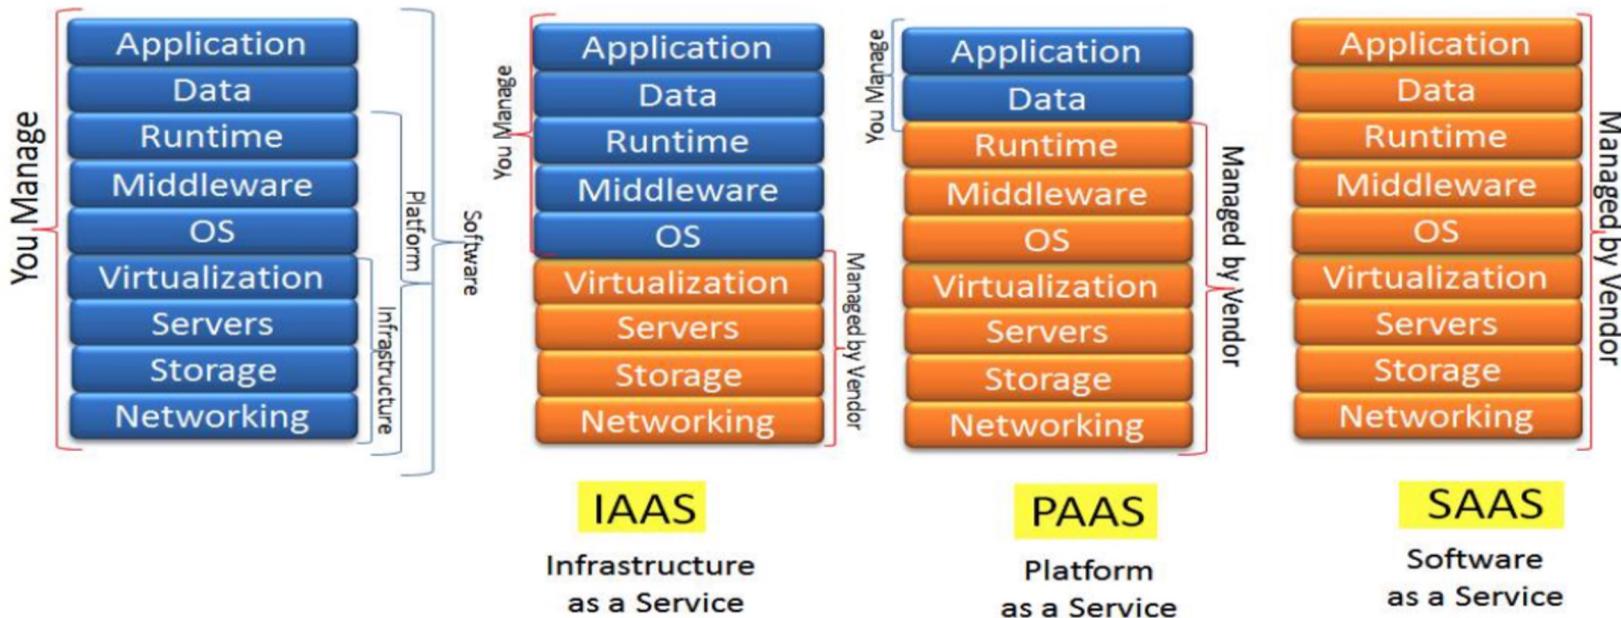
[Download CSV Summary](#)

Instance Name	VCPUs	Disk	RAM	Uptime
dminer-x86-ngz	1	10	2GB	1 week
dminer-x86-1	1	10	2GB	1 week
dminer-11.2-x86	1	10	2GB	1 week
dminer-x86-2	1	10	2GB	1 week
jibutler-x86	8	40	4GB	3 days, 2 hours

OpenStack IaaS services view



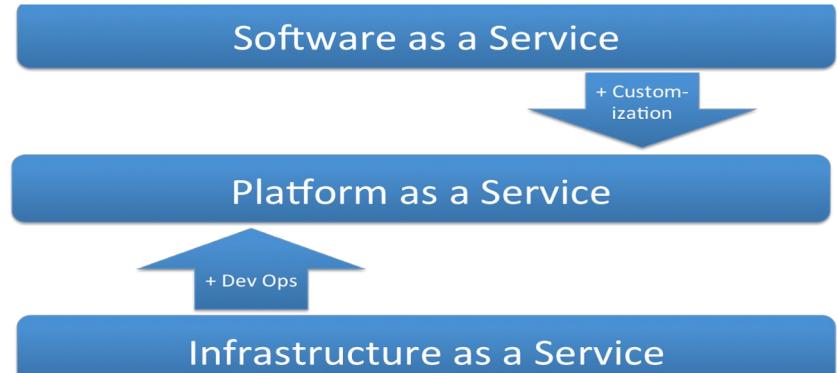
Platform As A Service



Taken and Slightly changed from
<https://cloudcelebrity.wordpress.com/2011/11/22/introduction-to-cloud-services-iaas-paas-saas/>

PaaS main concepts

- A platform is anything you can leverage to accomplish something in a simpler, faster, or otherwise better way than you could without.
- As a programmer, you leverage pre-existing code rather than starting from scratch and writing everything.
- The most well-known software platforms for desktop software are Windows and Mac OS



Platform as a Service

PaaS provides a framework for:

- **Application development:** A PaaS platform either provides the means to use programs you create in a supported language or offers a visual development environment that writes the code for you.
- **Collaboration:** Many PaaS systems are set up to allow multiple individuals to work on the same projects.
- **Data management:** Tools are provided for accessing and using data in a data store. **Instrumentation, performance, and testing:** Tools are available for measuring your applications and optimizing their performance.

Platform as a Service

PaaS provides a framework for:

- **Storage:** Data can be stored in either the PaaS vendor's service or accessed from a third- party storage service.
- **Transaction management:** Many PaaS systems provide services such as transaction managers or brokerage service for maintaining transaction integrity.

PaaS systems exist to allow you to create software that can be hosted as SaaS systems

Platform as a Service

PaaS cloud are designed to allow in developing robust, scalable, and hopefully portable applications.

On this list would be the following attributes:

- **Separate** of data management from the user interface | Reliance on cloud computing standards
- An integrated development environment (**IDE**)
- **Lifecycle** management tools
- **Multi-tenant** architecture support, security, and scalability
- Performance **monitoring**, testing, and optimization tools

A PaaS provides the tools needed to construct different types of applications that can work together in the same environment (e.g. mashup of multiple data sources).

In practice...

- To make it easier for you to run your website or web application no matter how much traffic it gets.
- You just deploy your application and the service figures out what to do with it.
- A platform as a service should handle scaling seamlessly for you so you can just focus on your website and the code running it.

PaaS examples

Google App Engine

OpenShift

Heroku

Force.com

Windows Azure

AWS Elastic Beanstalk

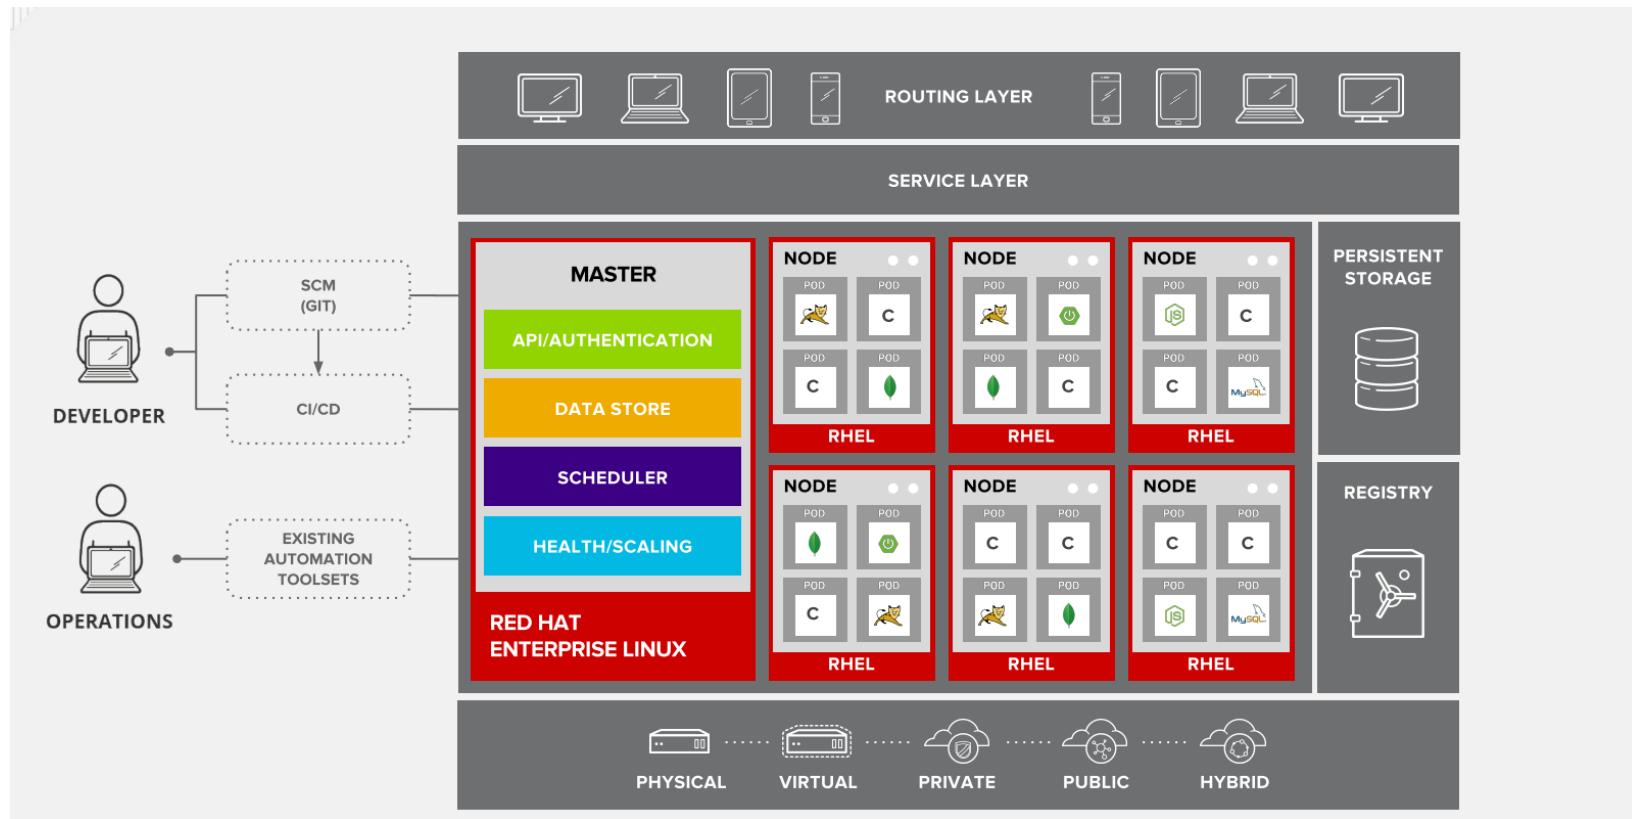
Apache Stratos

Magento Commerce

AWS Lamda

SAP Cloud

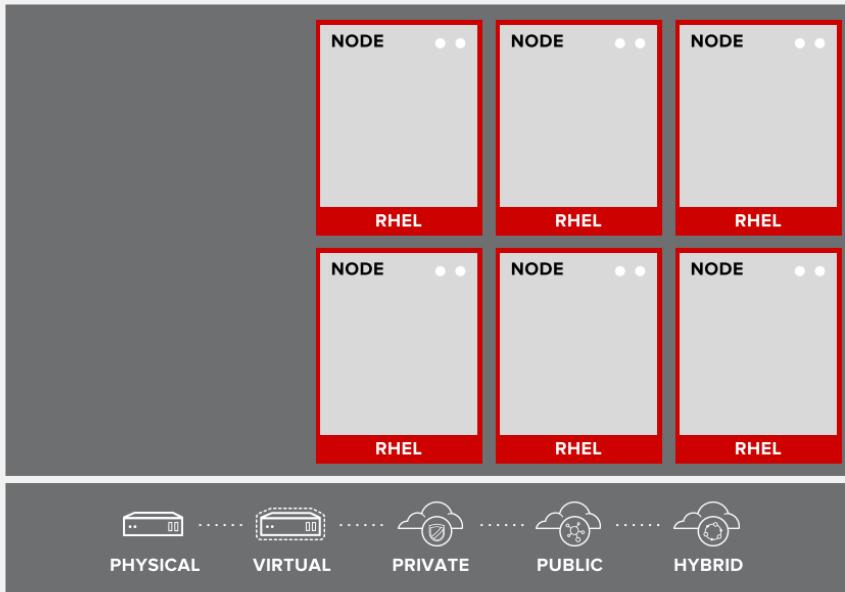
OpenShift Architecture



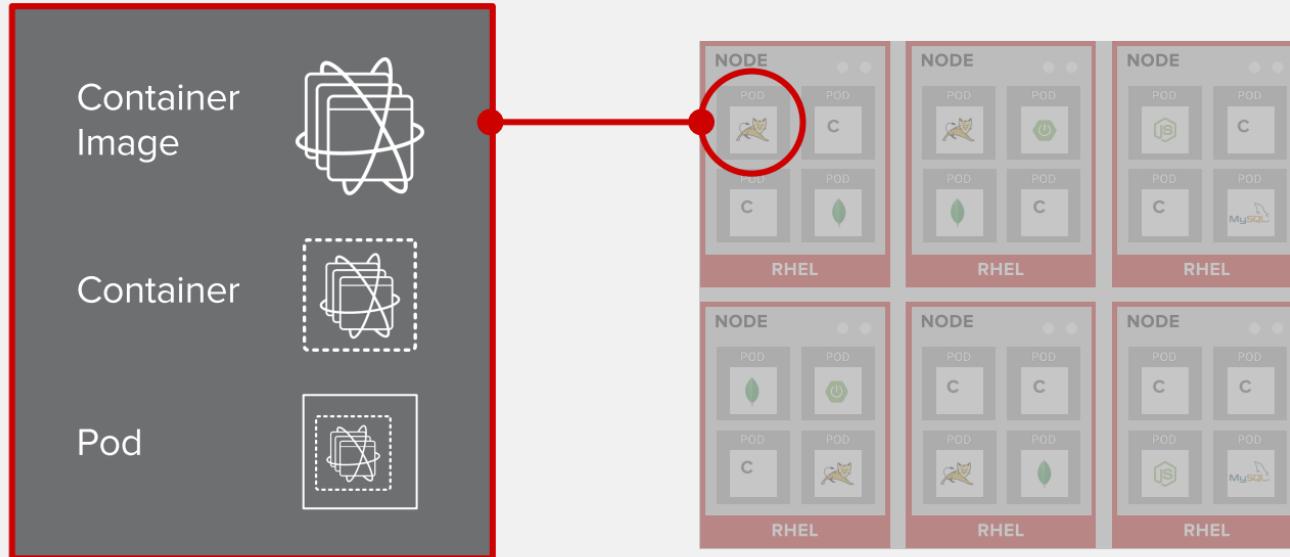
You can choose an infrastructure



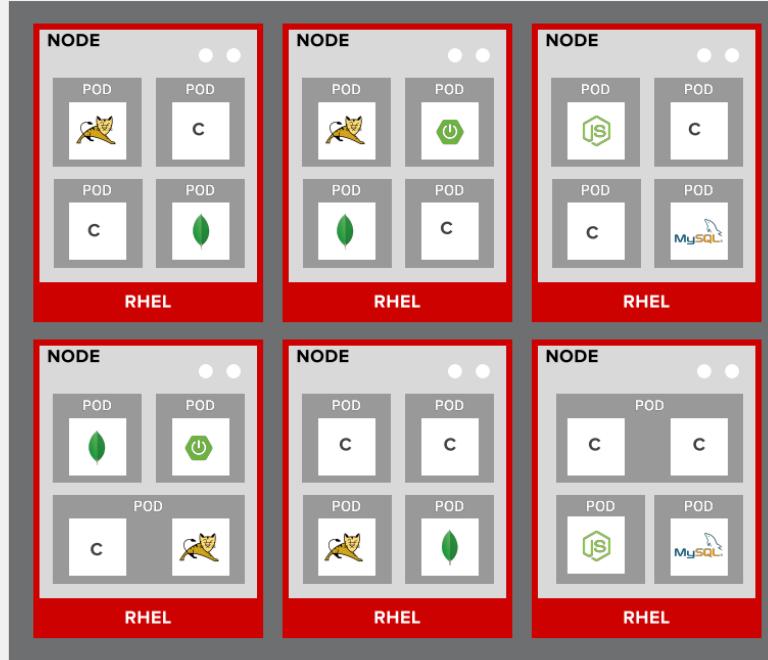
Nodes| RHEL instance where apps run



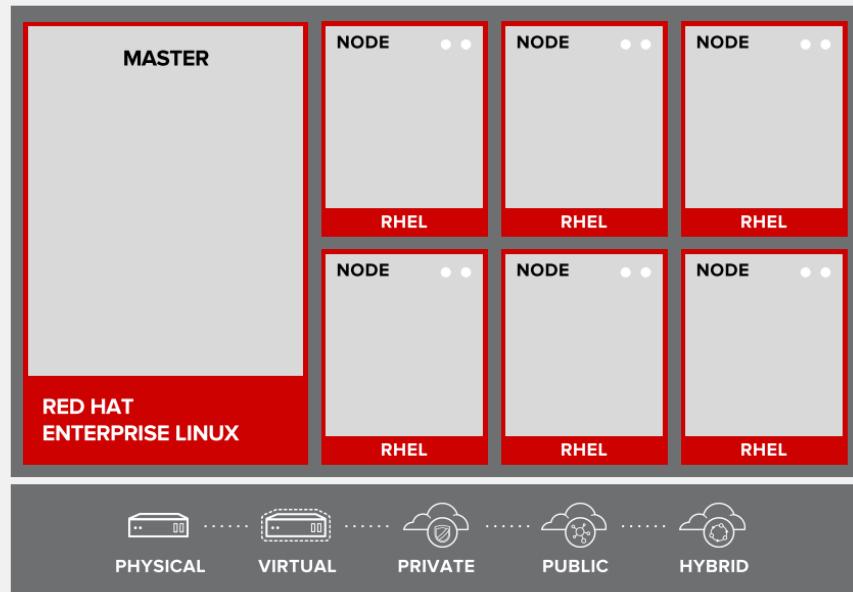
Apps run in containers



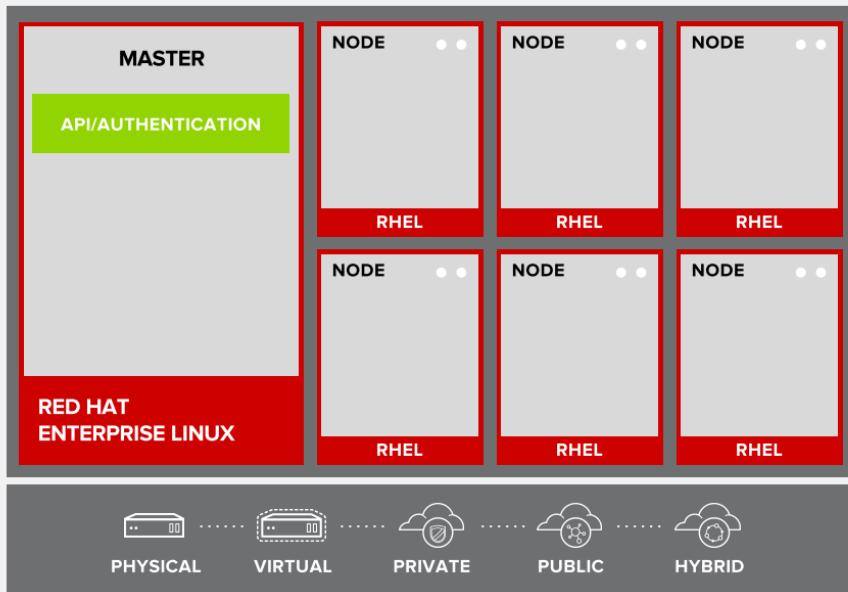
PODs are units of orchestration



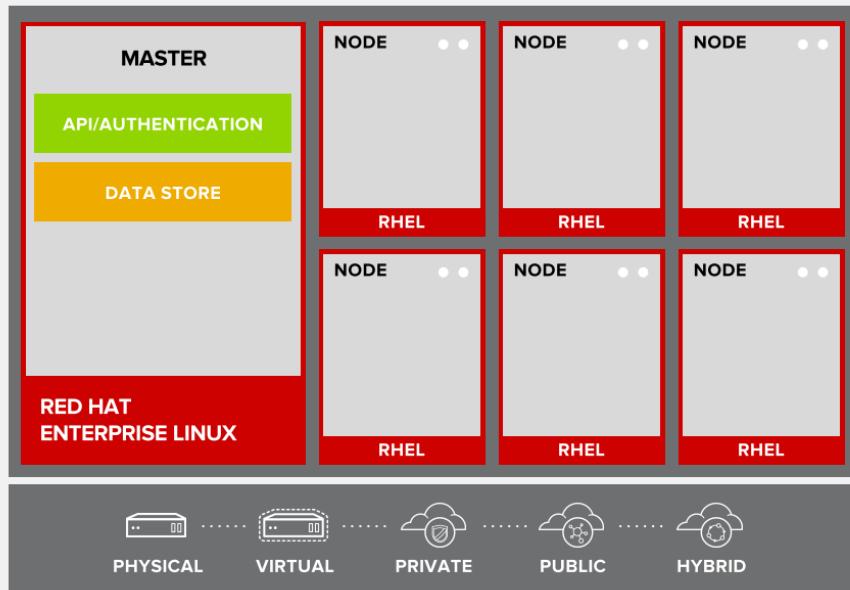
Masters are control plane



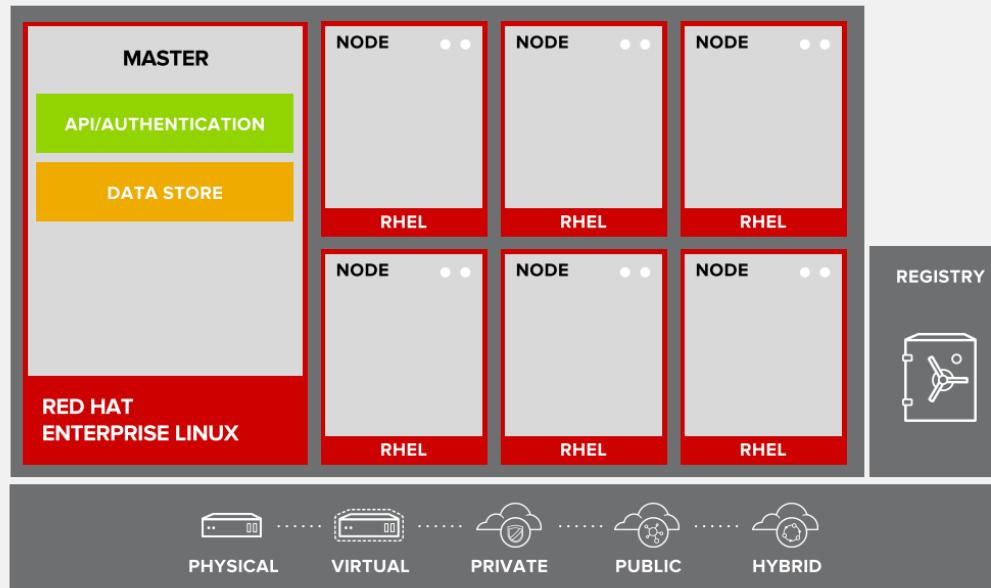
API AND AUTHENTICATION



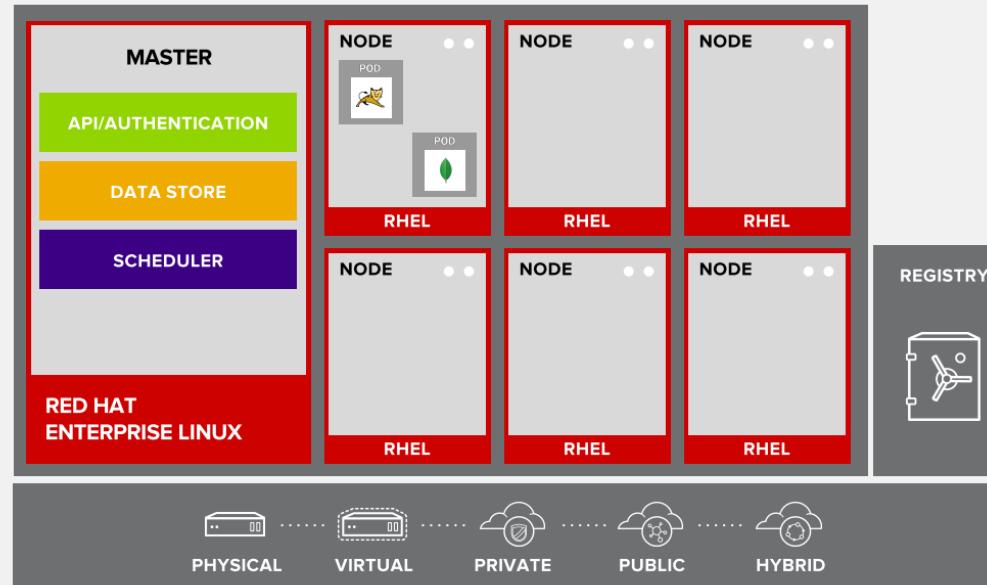
DESIRED AND CURRENT STATE



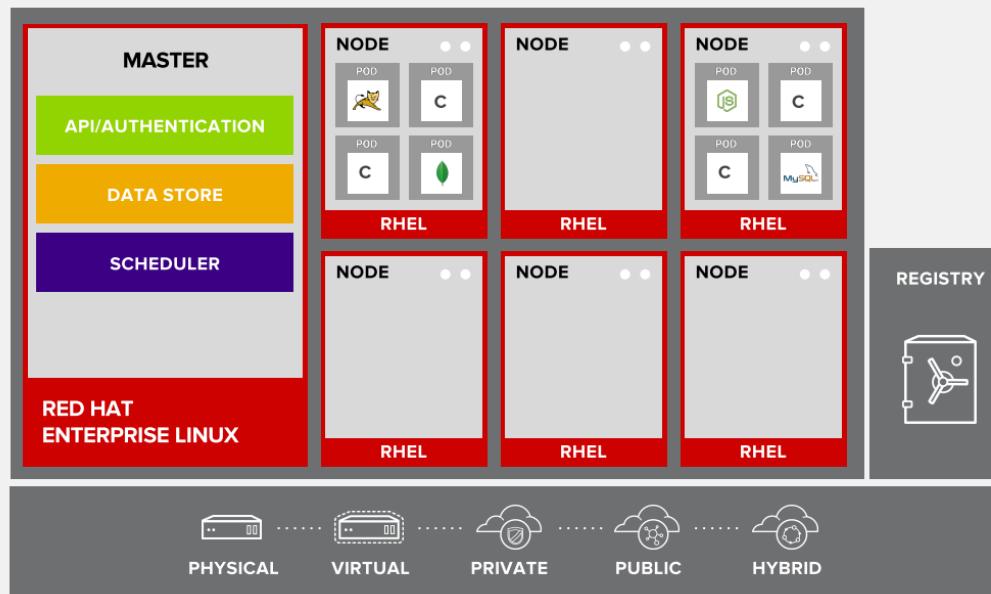
INTEGRATED CONTAINER REGISTRY



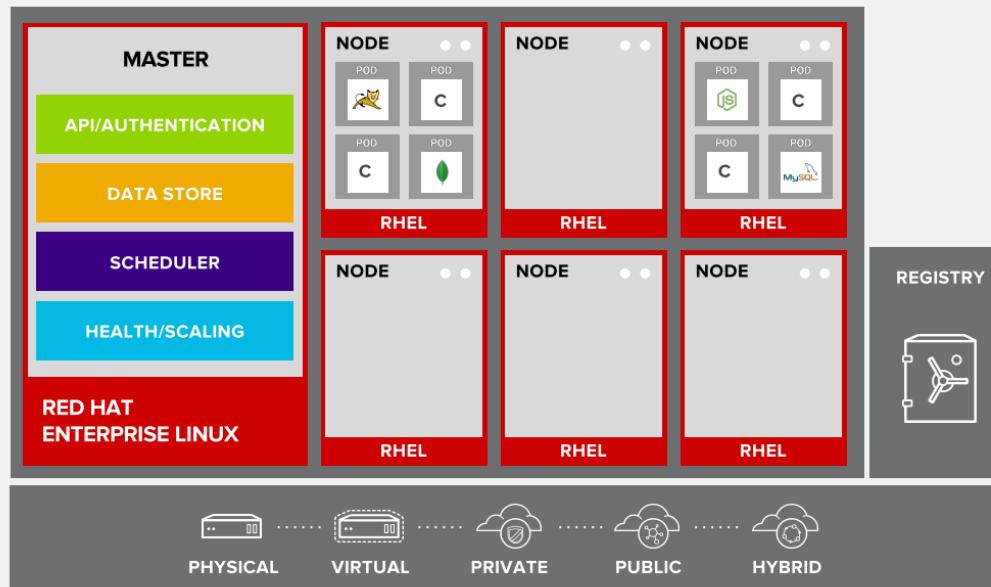
ORCHESTRATION AND SCHEDULING



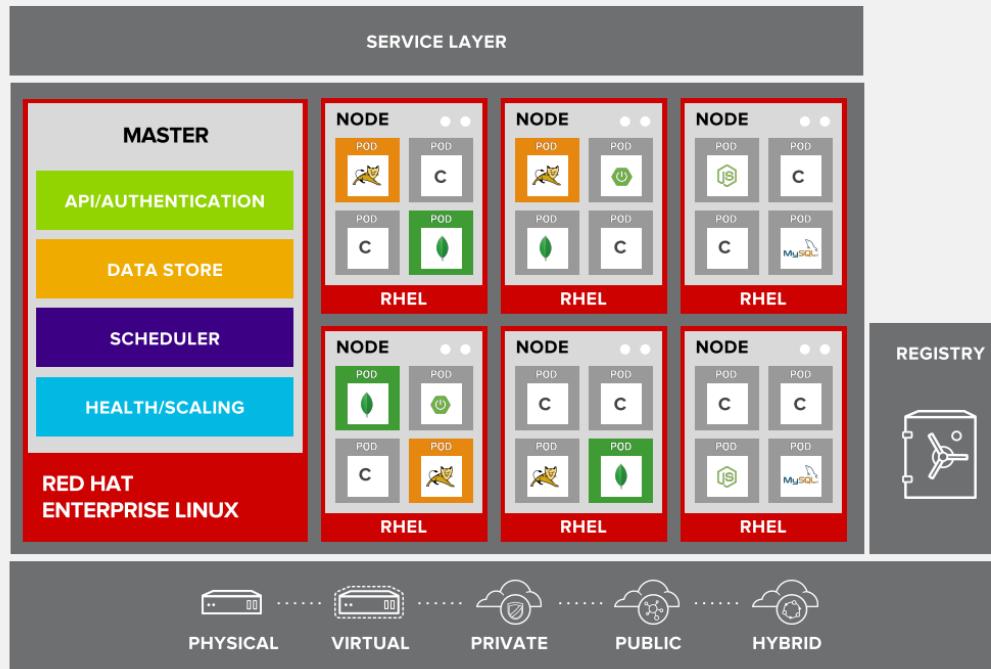
PLACEMENT BY POLICY



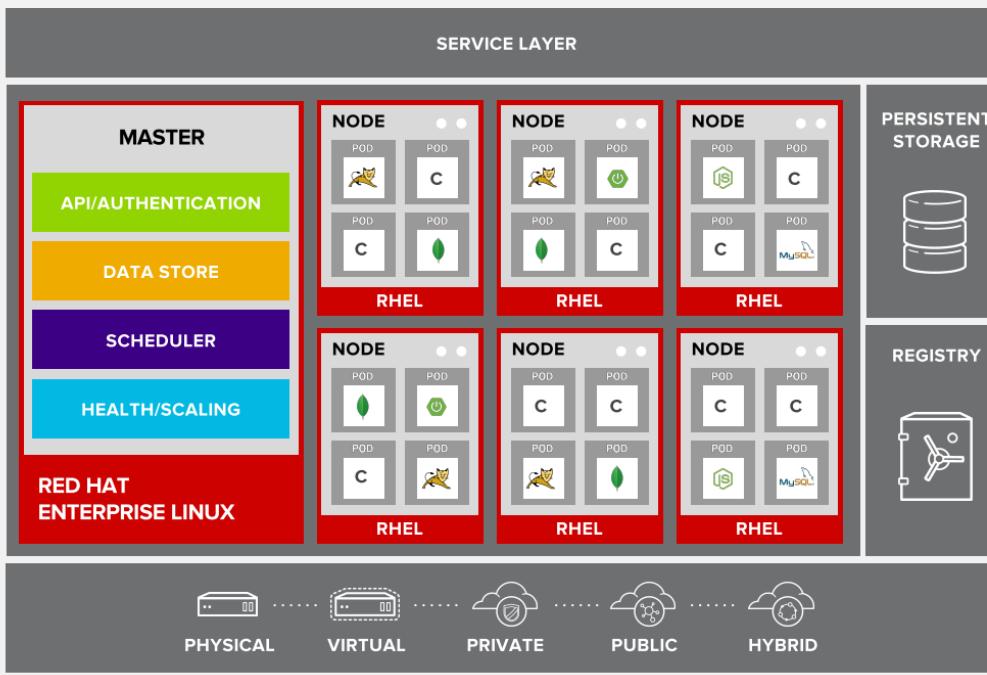
AUTOSCALING PODS



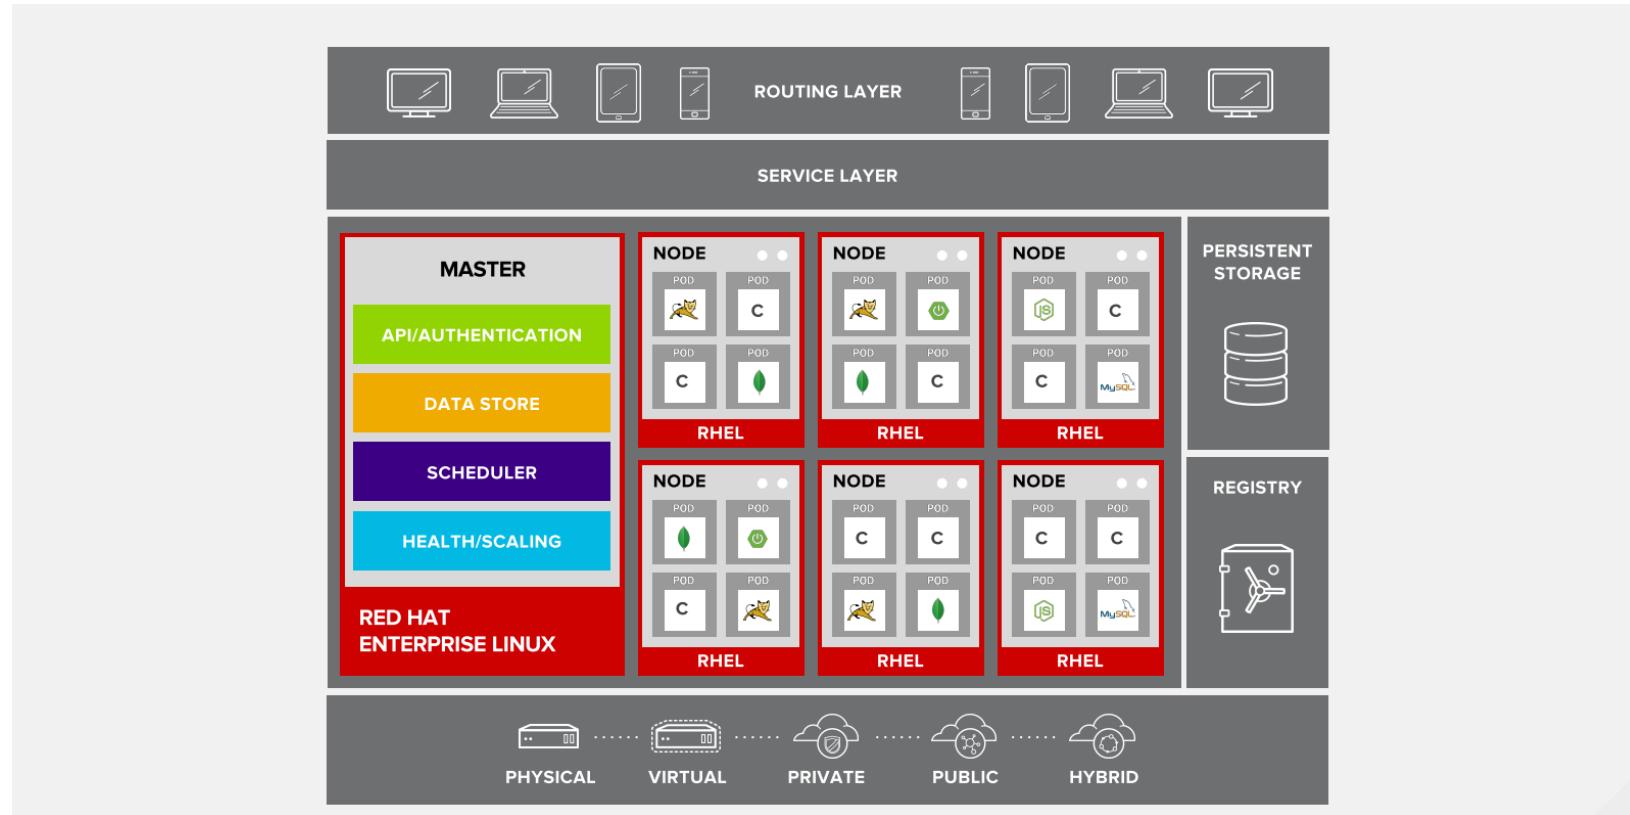
SERVICE DISCOVERY



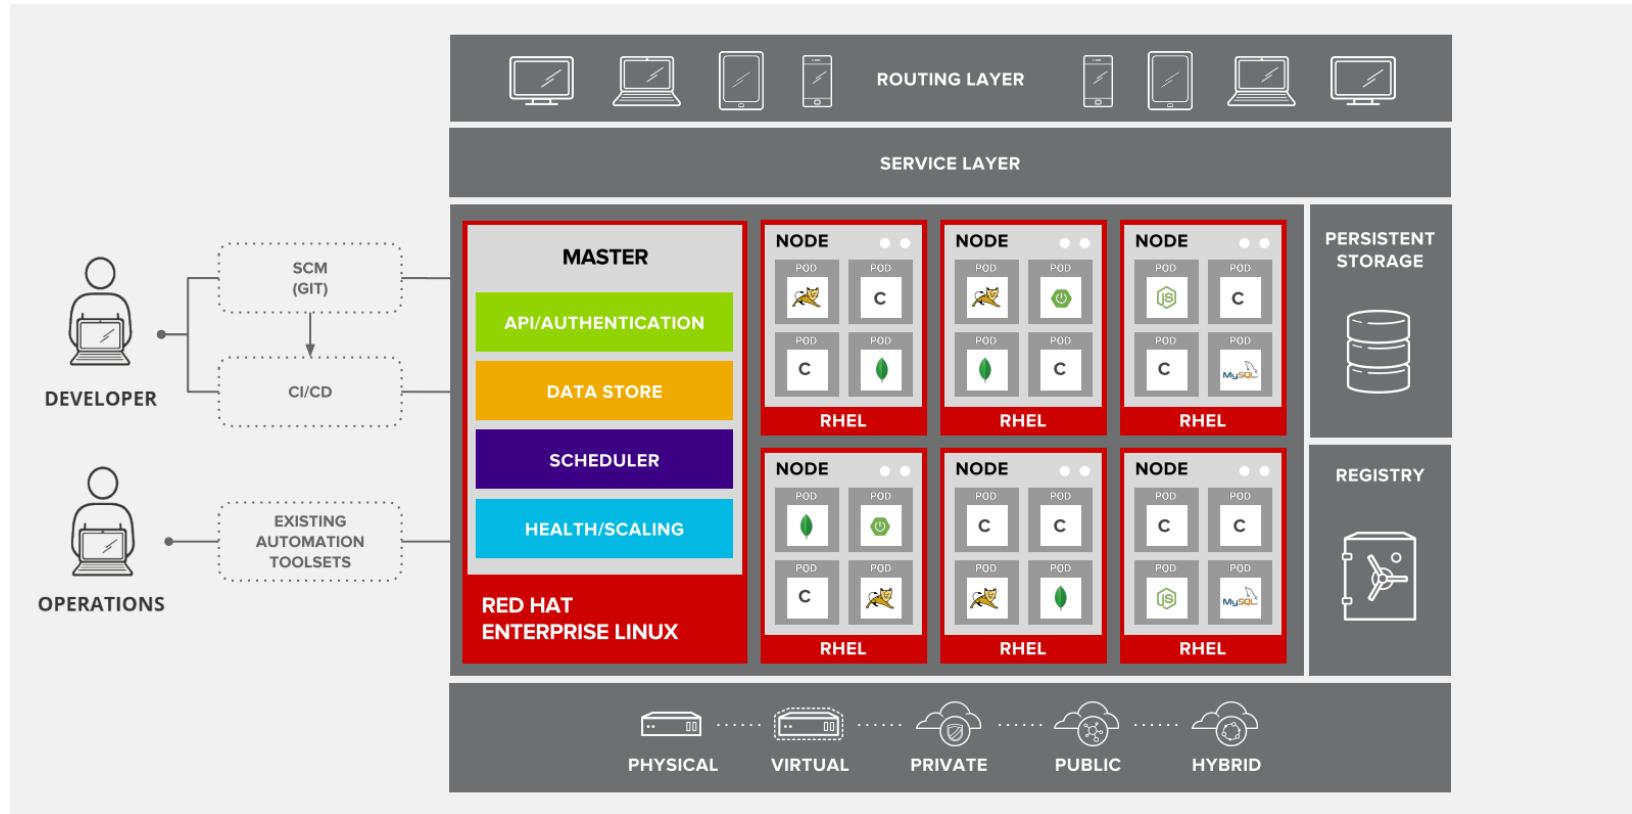
PERSISTENT DATA IN CONTAINERS



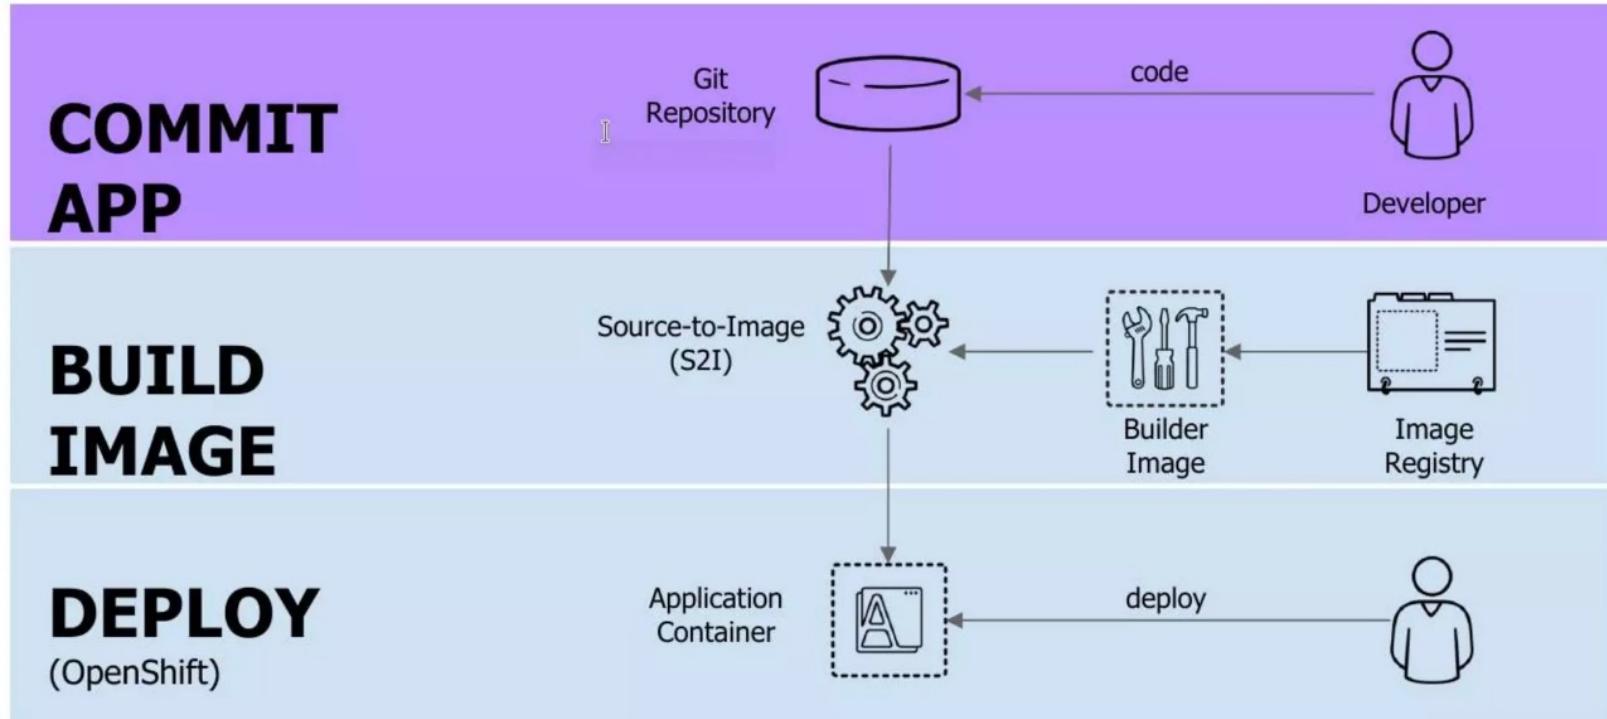
ROUTING AND LOAD-BALANCING



ACCESS VIA WEB, CLI, IDE AND API



How it works



console

The screenshot shows the Red Hat OpenShift Container Platform console interface. The top navigation bar includes the Red Hat logo, 'Red Hat OpenShift Container Platform', and a user dropdown for 'kubeadmin'. A blue banner at the top right states, 'You are logged in as a temporary administrative user. Update the cluster OAuth configuration to allow others to log in.' Below this, the project is set to 'serverless-example' and the application to 'all applications'. A 'View shortcuts' button and a three-dot menu are also present.

The left sidebar has a 'Developer' section with '+Add' and tabs for 'Topology' (which is selected), 'Monitoring', and 'Search'. Other sections include 'Builds', 'Pipelines', 'Helm', 'Project', 'Secrets', and 'Pods'.

The main content area is titled 'Topology' and displays a message: 'No workloads found. To add content to your project, create an application, component or service using one of these options.' Below this, there are eight cards arranged in two rows of four:

- From Git**: Import code from your git repository to be built and deployed.
- Container Image**: Deploy an existing image from an image registry or image stream tag.
- From Dockerfile**: Import your Dockerfile from your git repo to be built and deployed.
- YAML**: Create resources from their YAML or JSON definitions.

- From Catalog**: Browse the catalog to discover, deploy and connect to services.
- Database**: Browse the catalog to discover database services to add to your application.
- Operator Backed**: Browse the catalog to discover and deploy operator managed services.
- Helm Chart**: Browse the catalog to discover and install Helm Charts.

that's all, have fun

"So long
and thanks
for all the fish"